

ECE 699: Lecture 6

AXI Interfacing

Using DMA & AXI4-Stream

Required Reading

The ZYNQ Book

- *Chapter 19: AXI Interfacing*
- *Chapter 2.3: Processing System – Programmable Logic Interfaces*
- *Chapter 9.3: Buses*
- *Chapter 10.1: Interfacing and Signals*

Required Reading

ARM AMBA AXI Protocol v1.0: Specification

- *Chapter 1: Introduction*
- *Chapter 2: Signal Descriptions*
- *Chapter 3: Channel Handshake*
- *Chapter 4: Addressing Options*
- *Chapter 9: Data Buses*

Recommended Reading

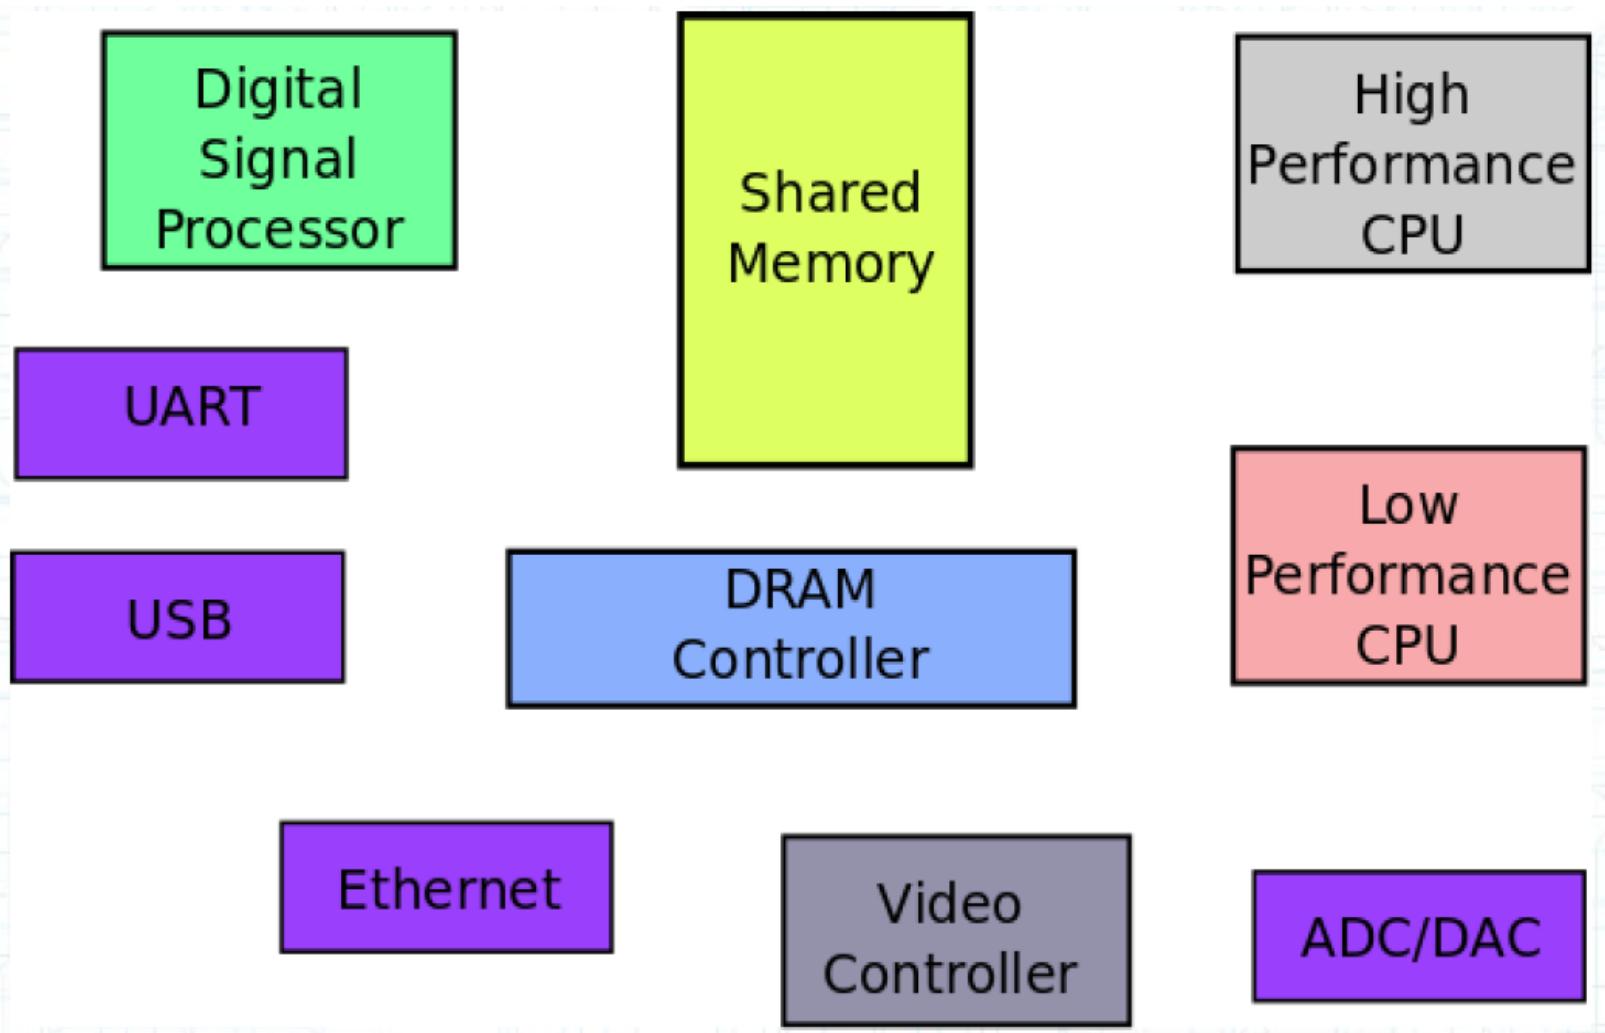
P. Schaumont, A Practical Introduction to Hardware/Software Codesign, 2nd Ed.

- *Chapter 10: On-Chip Buses*

M.S. Sadri, ZYNQ Training
(presentations and videos)

- *Lesson 1 : What is AXI?*
- *Lesson 2 : What is an AXI Interconnect?*
- *Lesson 3 : AXI Stream Interface*

Components of Today's Systems-on-Chip



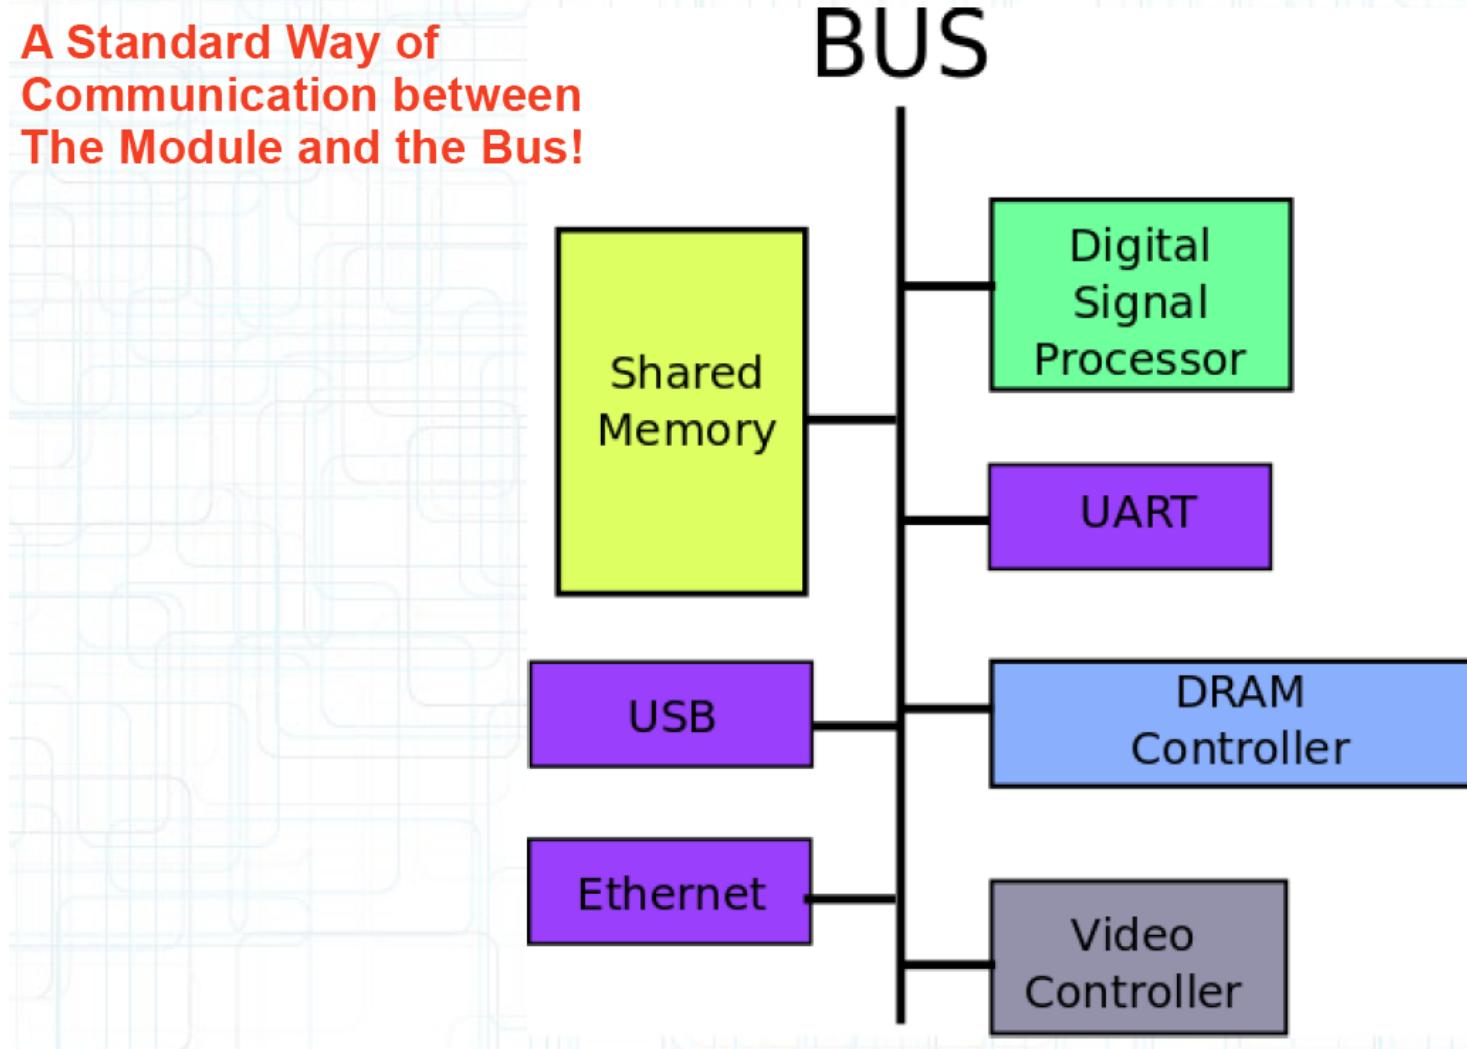
Source: M.S. Sadri, Zynq Training

Connectivity Requirements

- A standard
 - All units talk based on that standard
 - All units can talk easily to each other
- Maintenance
 - Design is easily maintained/updated, debugged
- Re-use
 - Units can be easily re-used in other designs

SoC Buses

A Standard Way of
Communication between
The Module and the Bus!



Source: M.S. Sadri, Zynq Training

Solution Adopted in ZYNQ

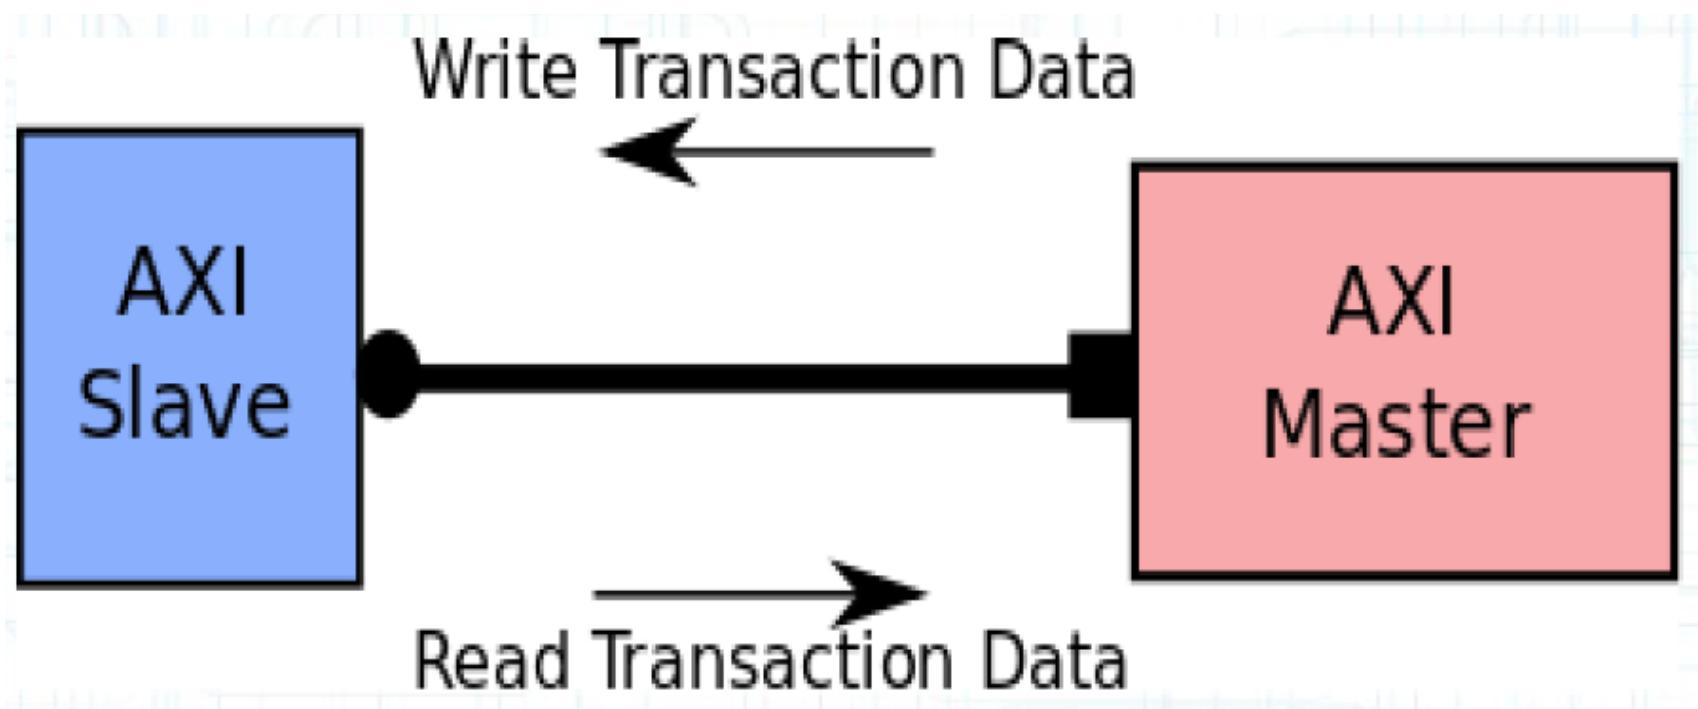
Advanced Microcontroller Bus Architecture (AMBA):
an open-standard, on-chip interconnect specification
for the connection and management of functional
blocks in system-on-a-chip (SoC) designs.
First version introduced by ARM in 1996.

AMBA Advanced eXtensible Interface 4 (AXI4):
the fourth generation of AMBA interface defined
in the AMBA 4 specification, targeted at
high performance, high clock frequency systems.
Introduced by ARM in 2010.

Basic Concepts

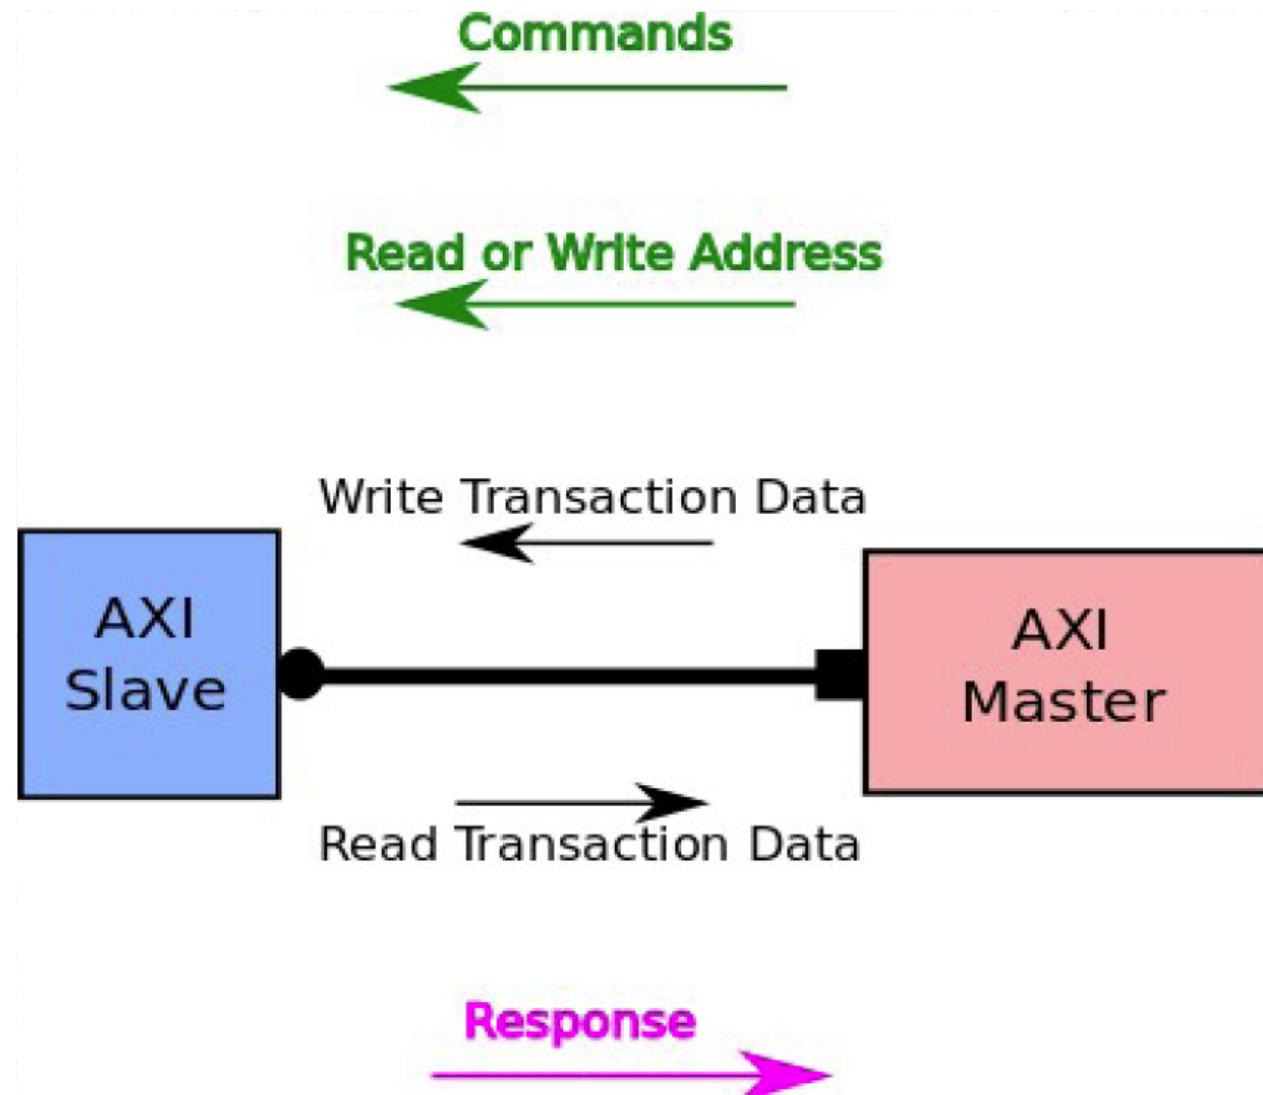
- Transaction :
 - Transfer of data from one point in the hardware to another point
- Master : Initiates the transaction
- Slave : Responds to the initiated transaction

Communication Between AXI Master and AXI Slave



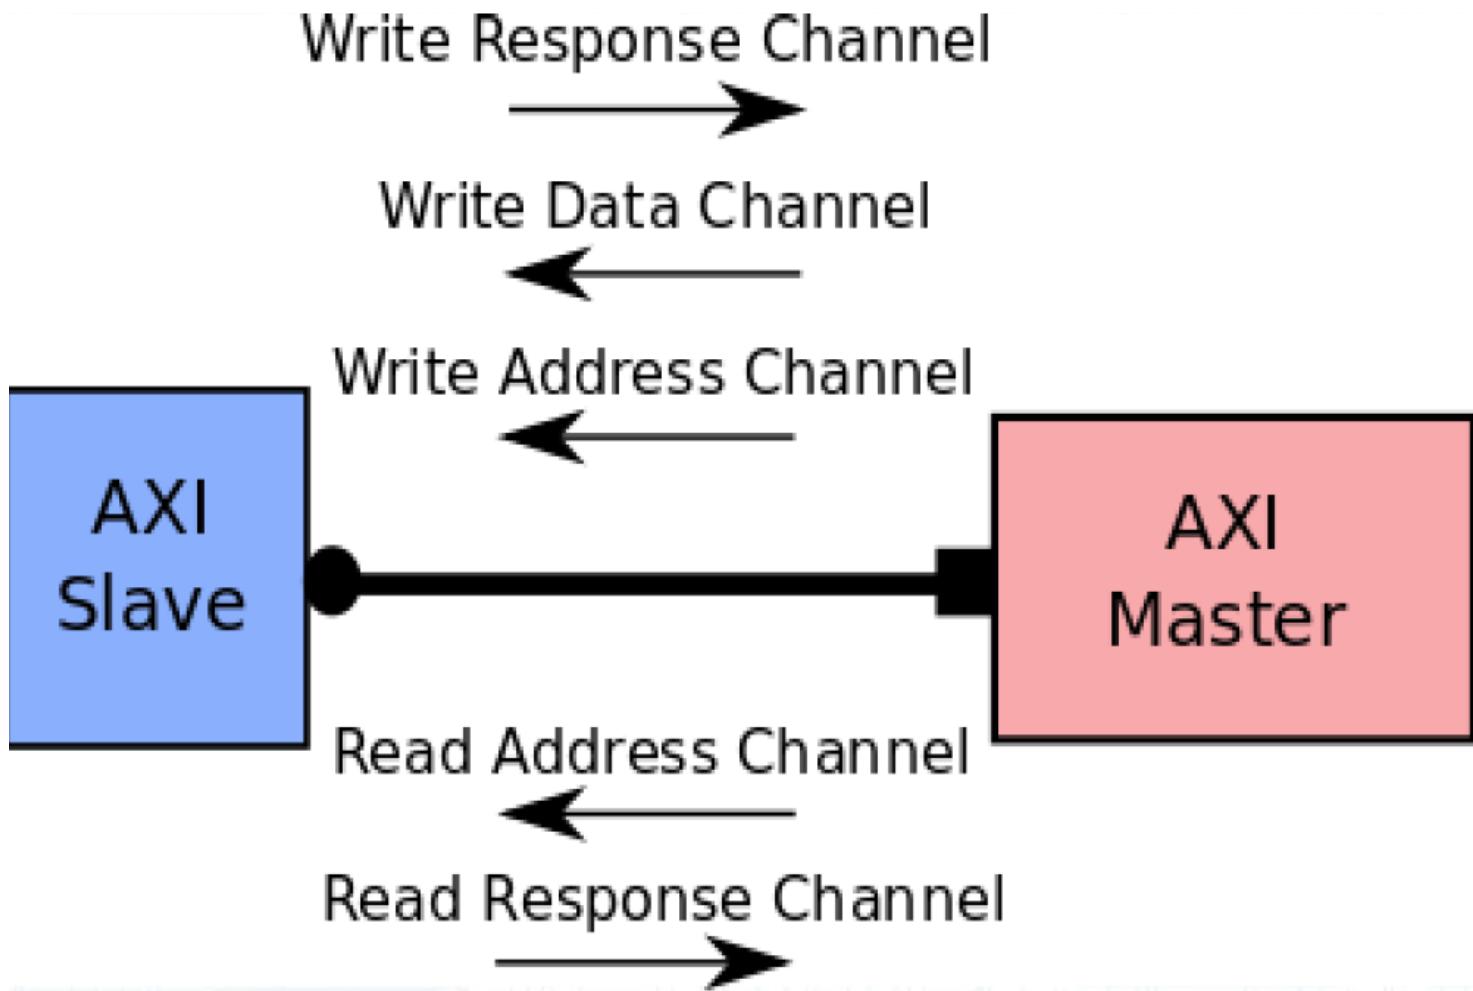
Source: M.S. Sadri, Zynq Training

Additional Information Exchanged Between AXI Master and AXI Slave



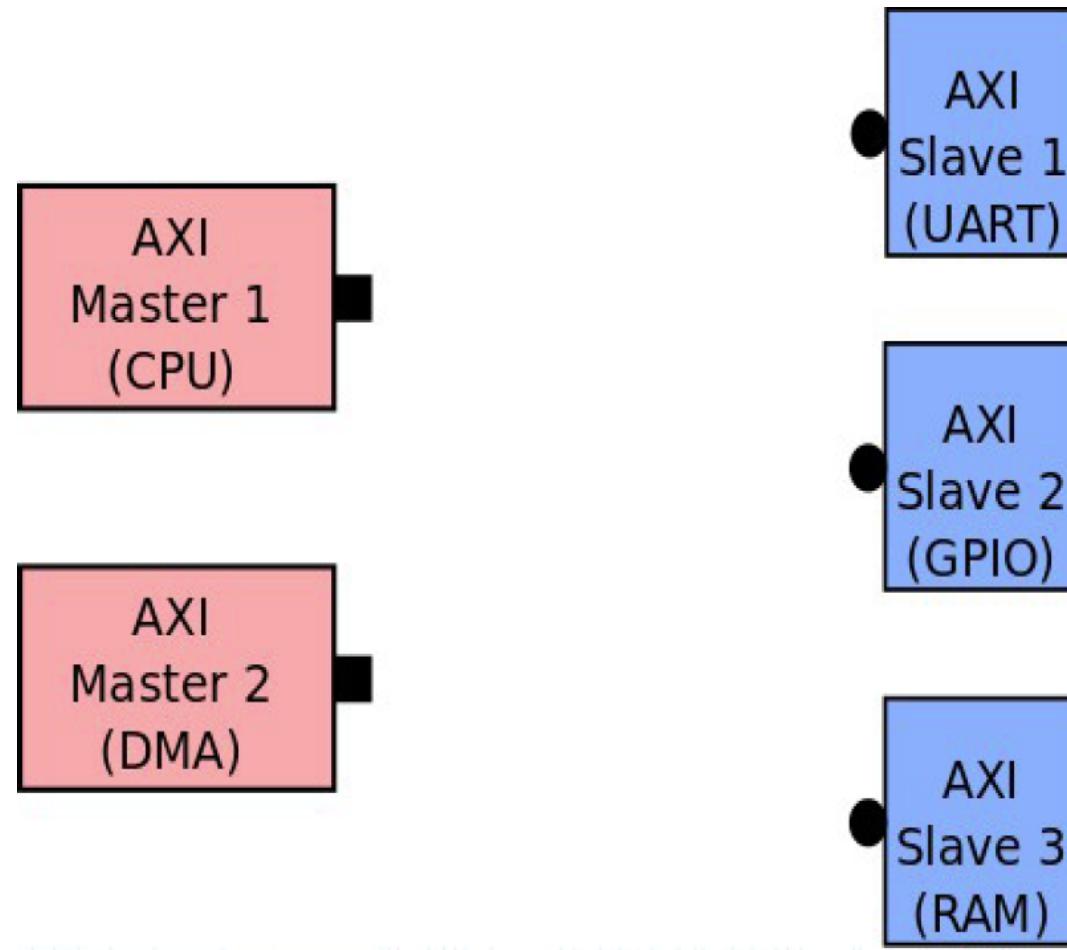
Source: M.S. Sadri, Zynq Training

Five Channels of AXI Interface



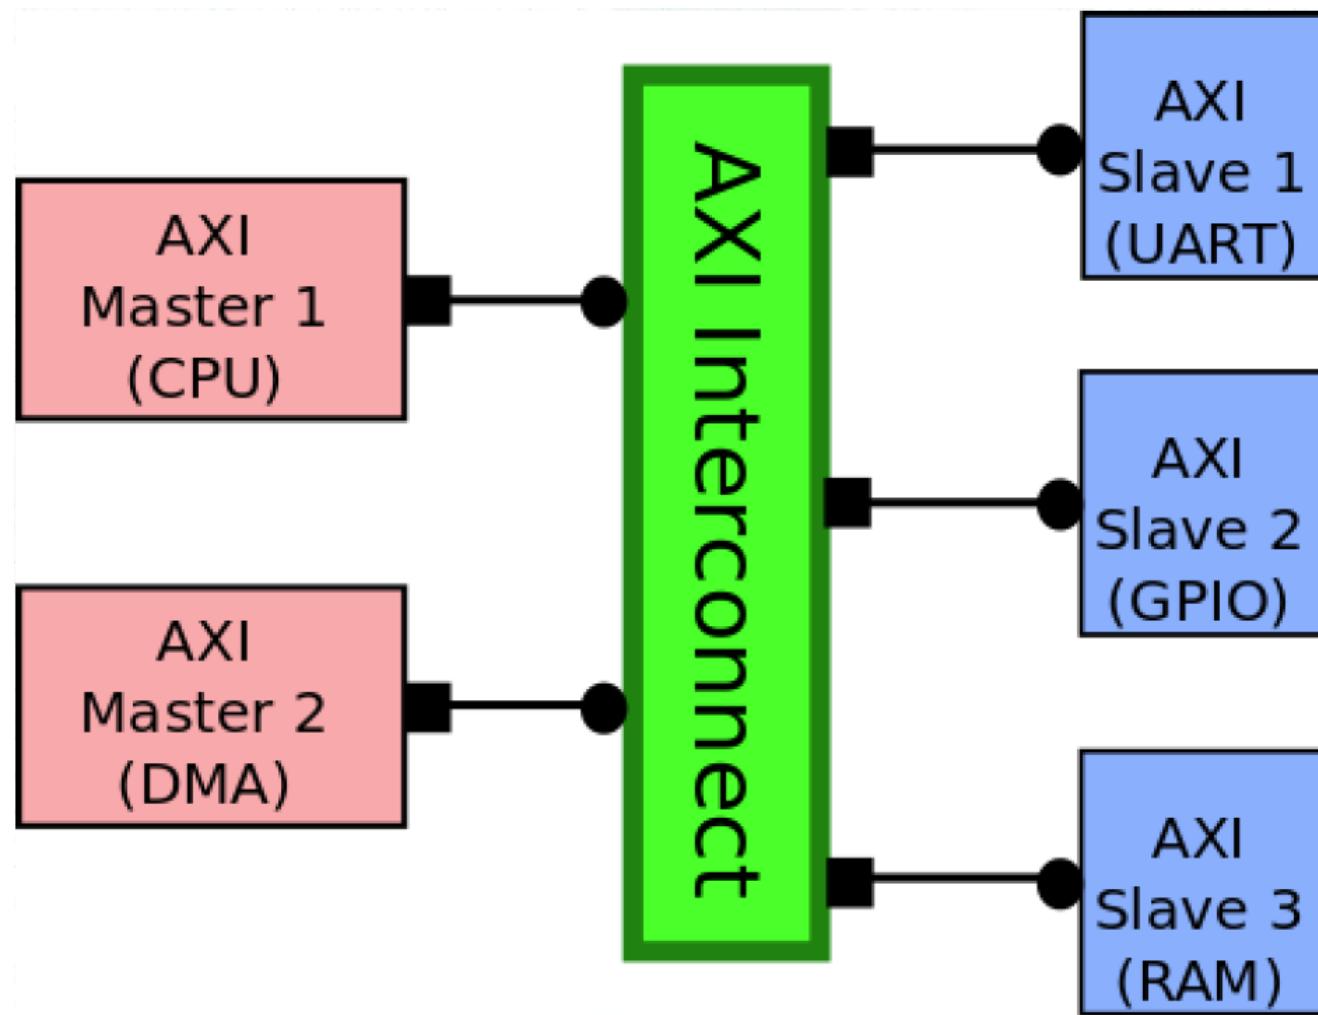
Source: M.S. Sadri, Zynq Training

Connecting Masters and Slaves

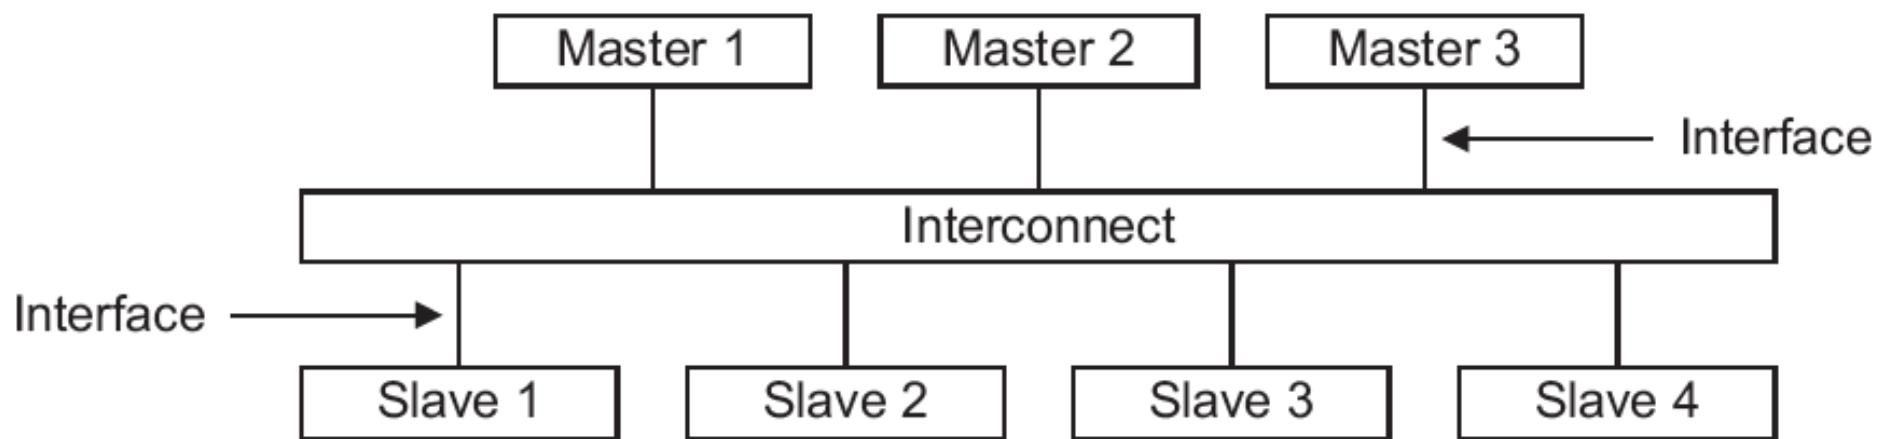


Source: M.S. Sadri, Zynq Training

AXI Interconnect



Interconnect vs. Interface



AXI Interfaces and Interconnects

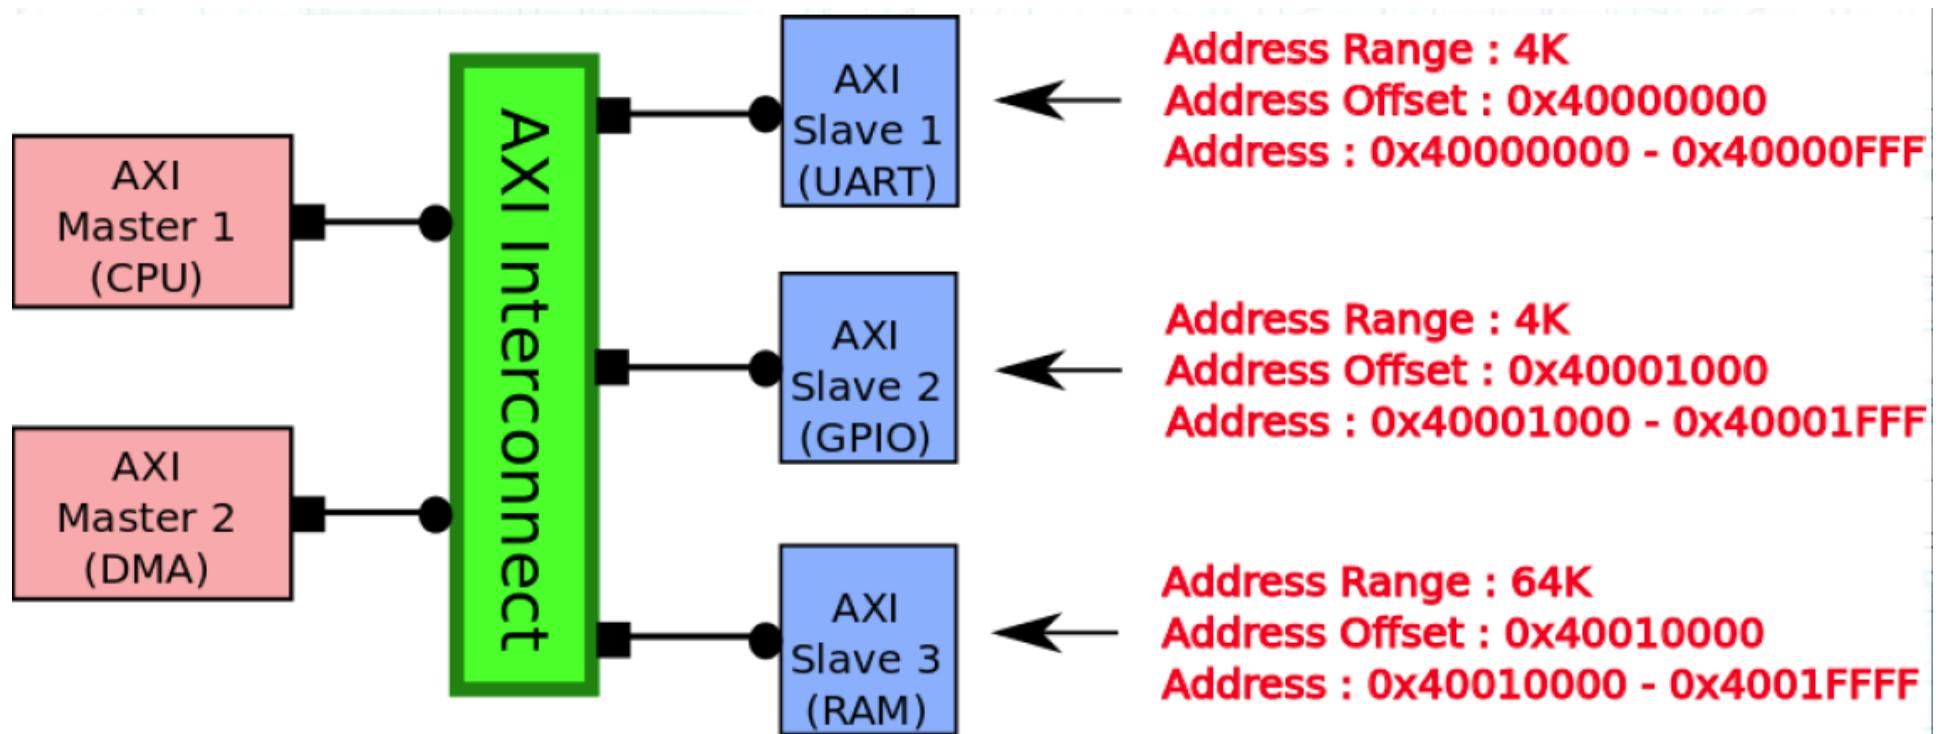
Interface

A point-to-point connection for passing data, addresses, and hand-shaking signals between master and slave clients within the system

Interconnect

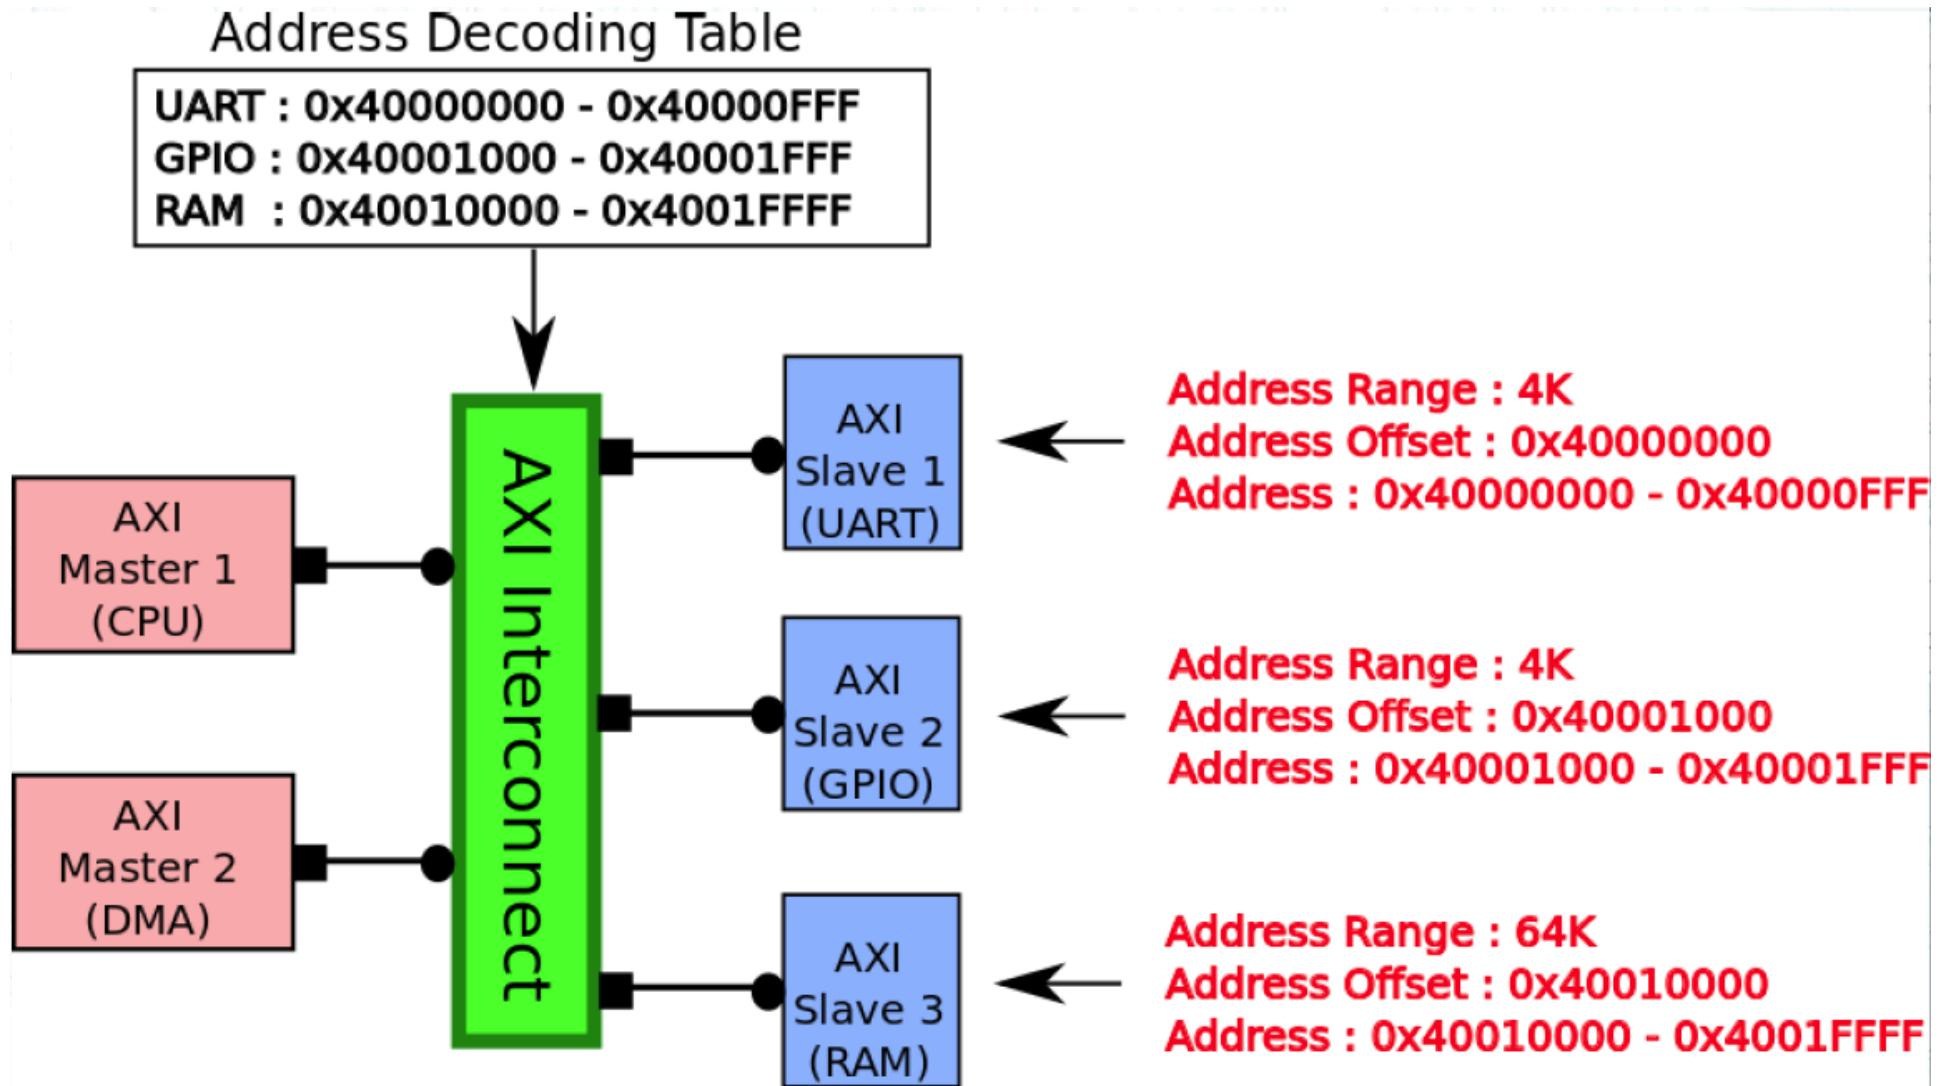
A switch which manages and directs traffic between attached AXI interfaces

Addressing of Slaves



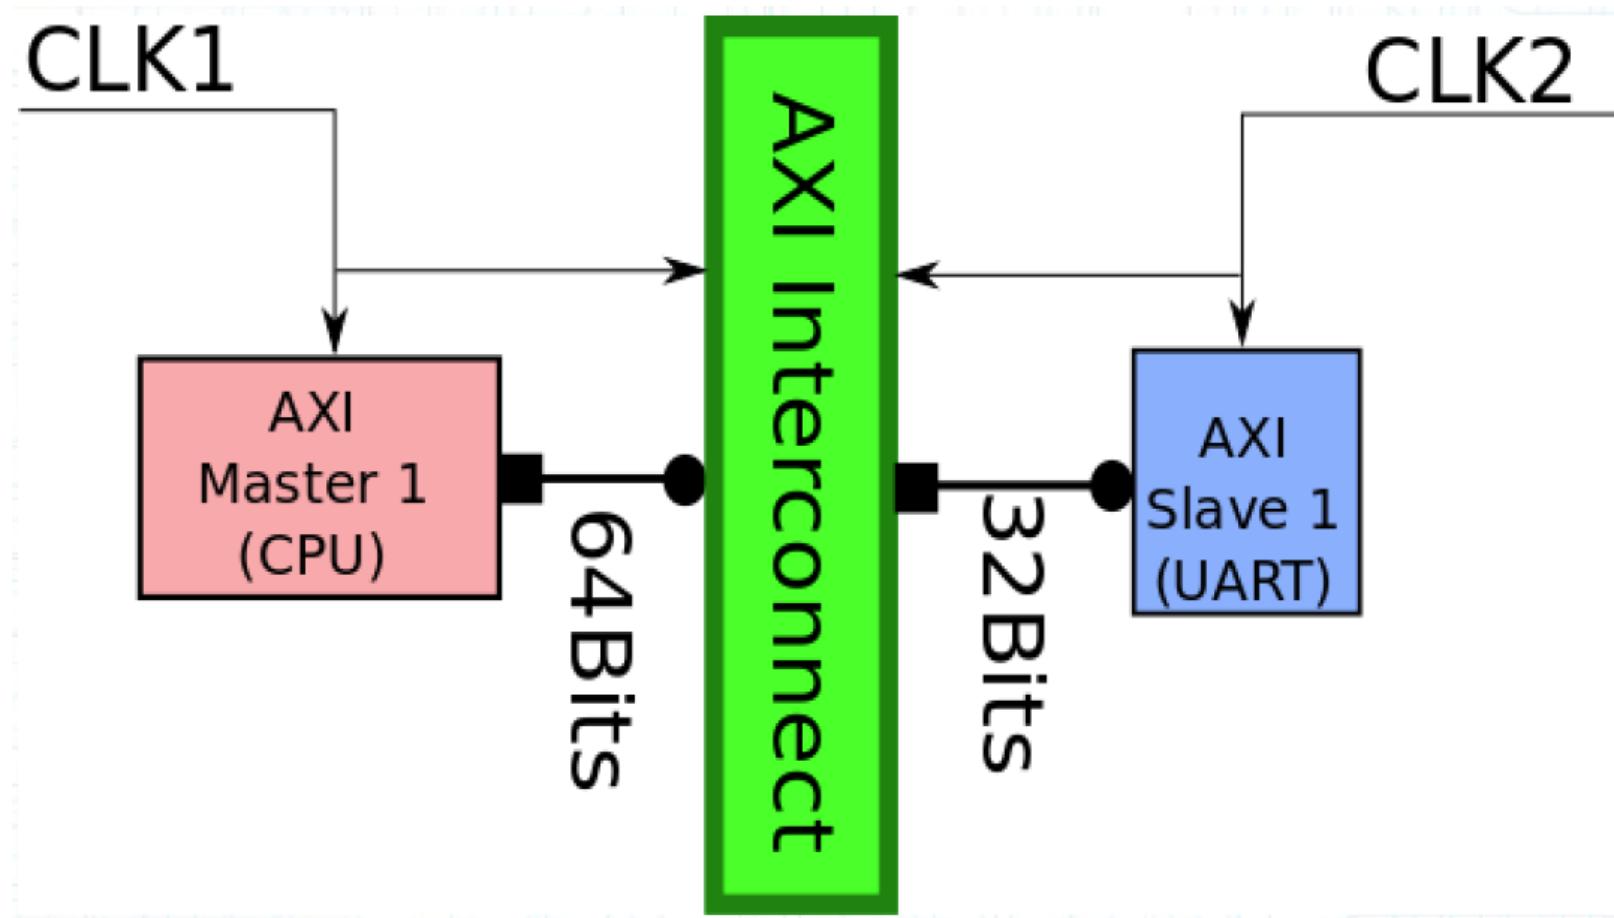
Source: M.S. Sadri, Zynq Training

AXI Interconnect Address Decoding



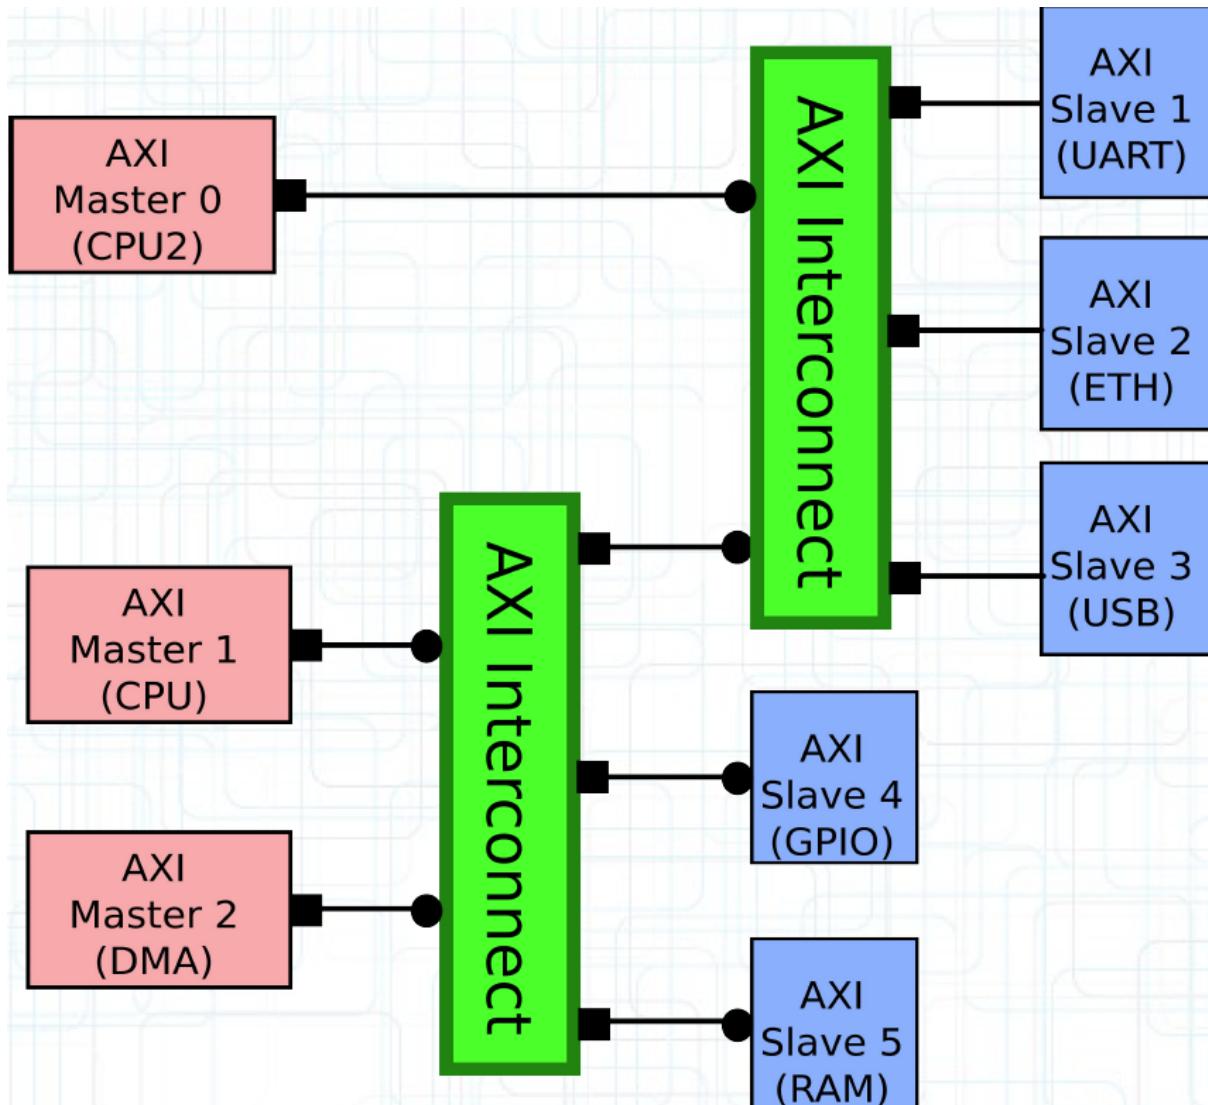
Source: M.S. Sadri, Zynq Training

Clock Domain and Width Conversion



Source: M.S. Sadri, Zynq Training

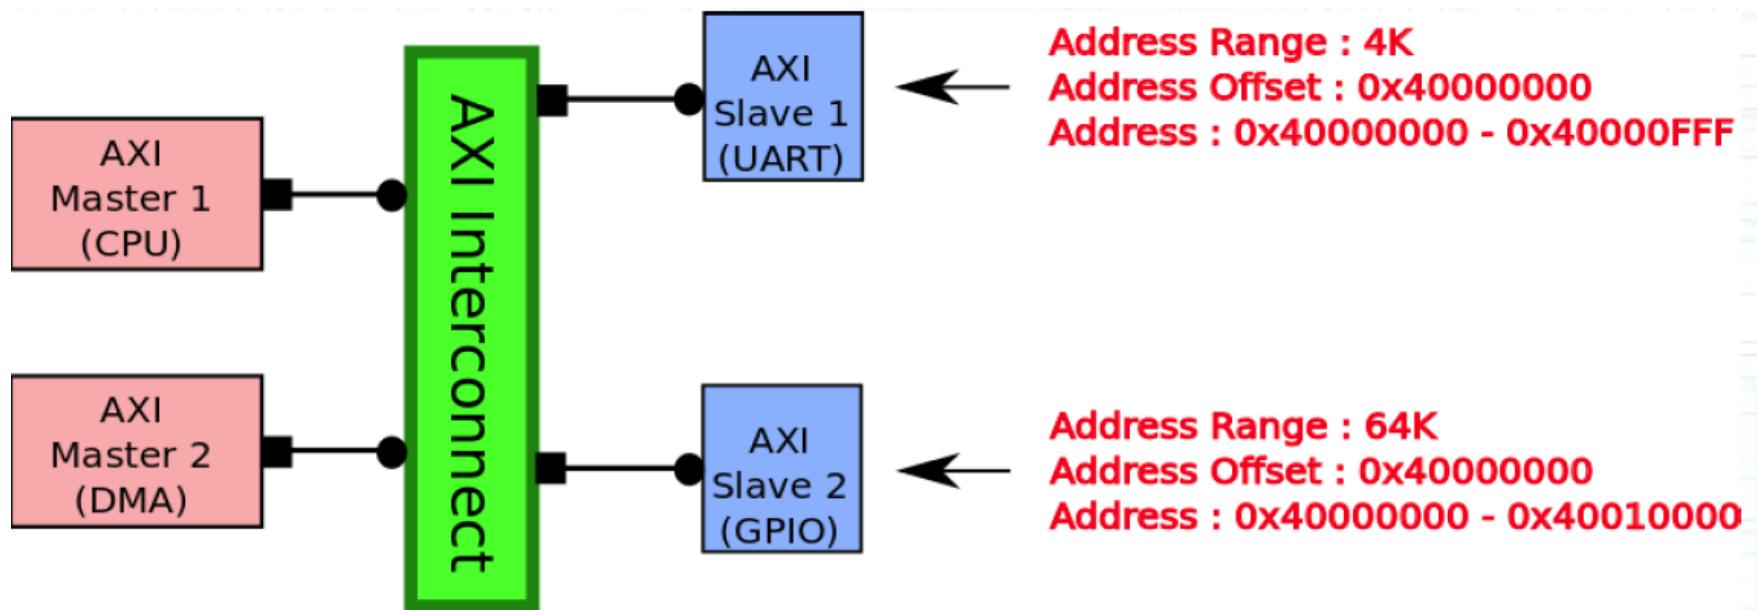
Hierarchical AXI Interconnects



Source: M.S. Sadri, Zynq Training

Simple Address Definition Rules

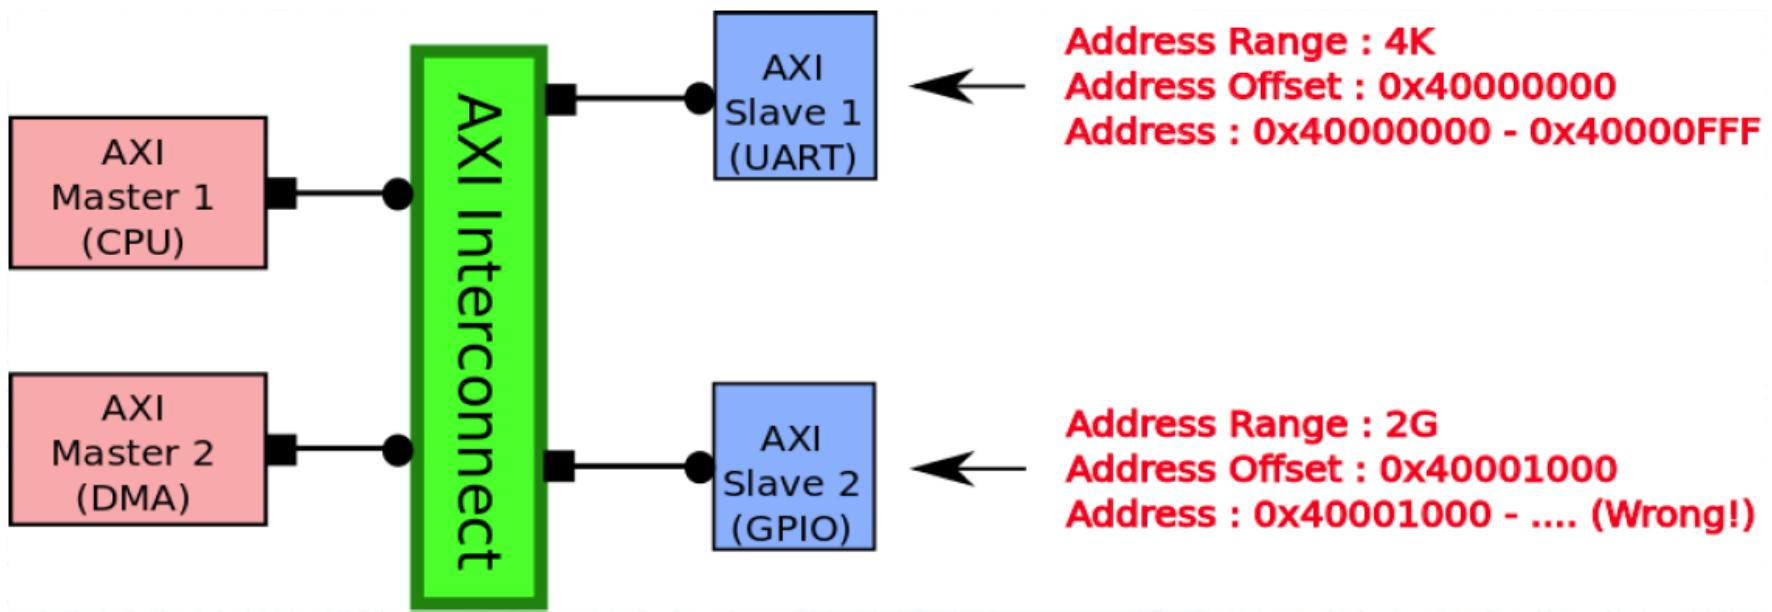
No Overlaps



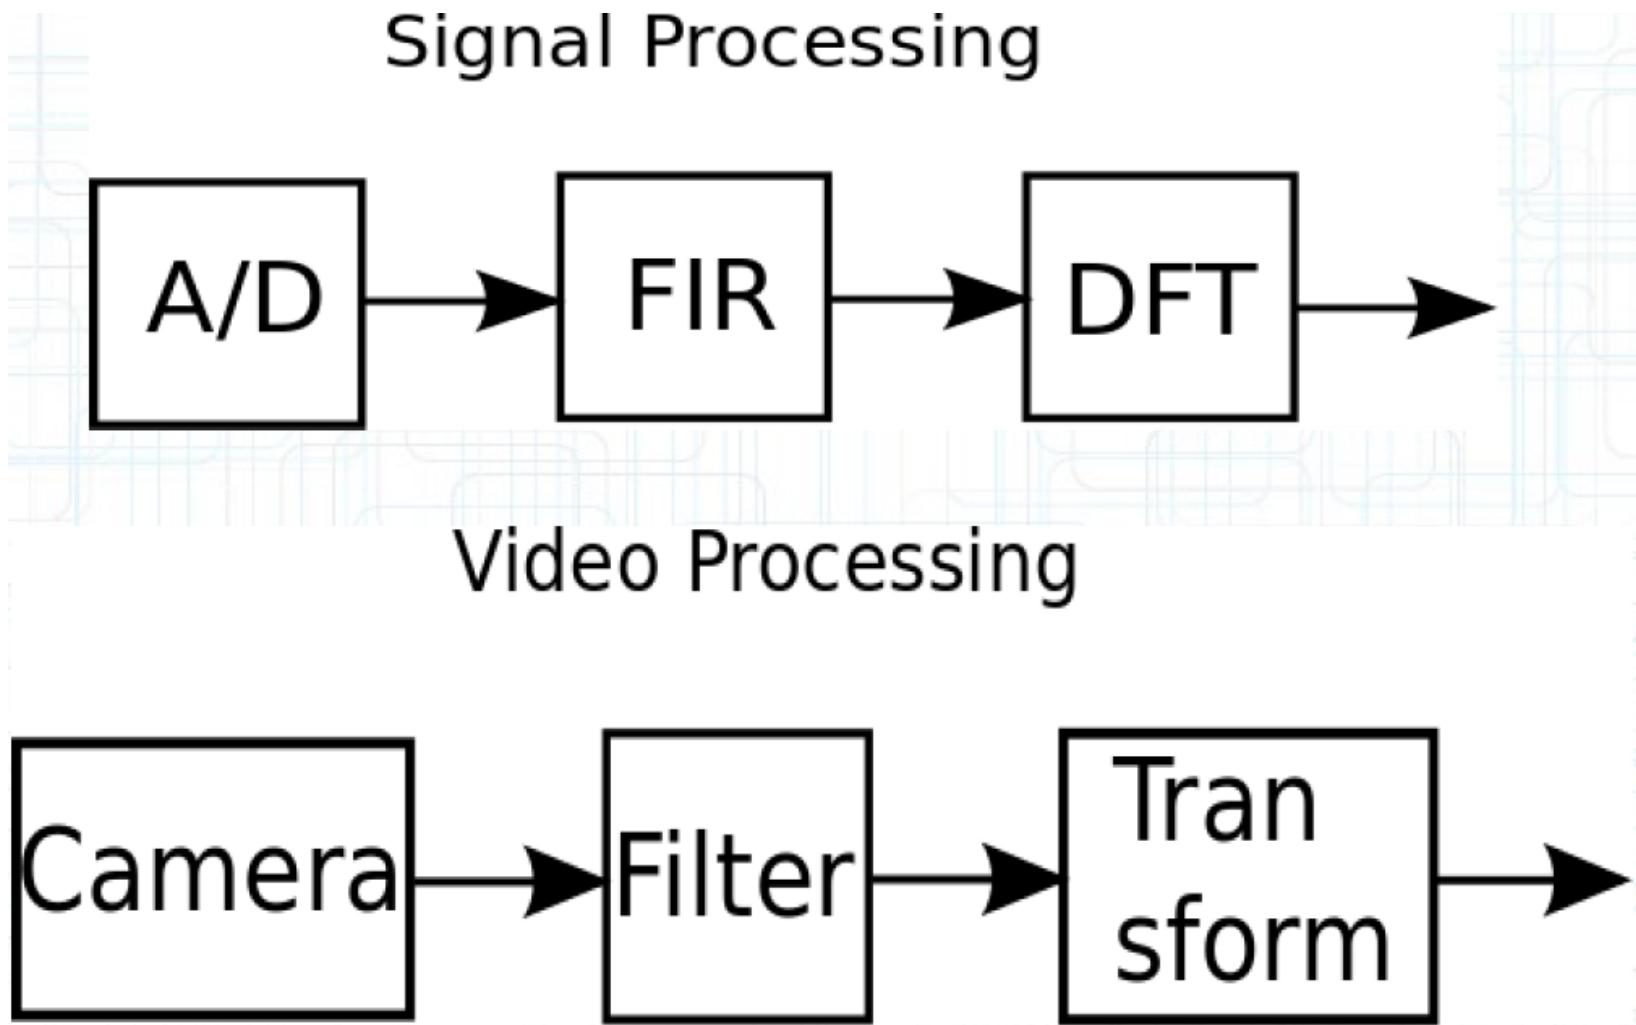
Source: M.S. Sadri, Zynq Training

Simple Address Definition Rules

Address Alignment

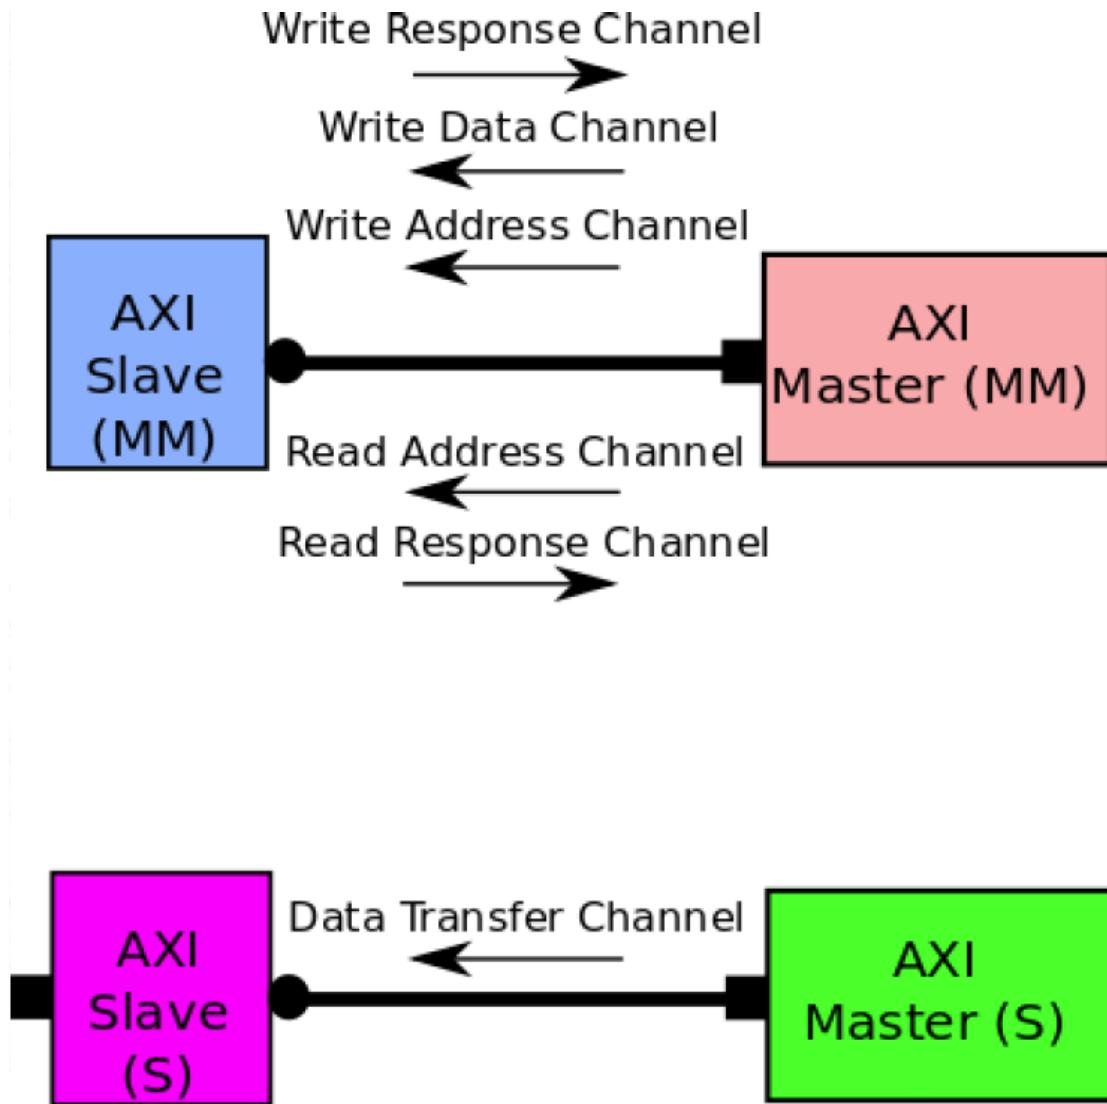


Point-to-Point Data Flows



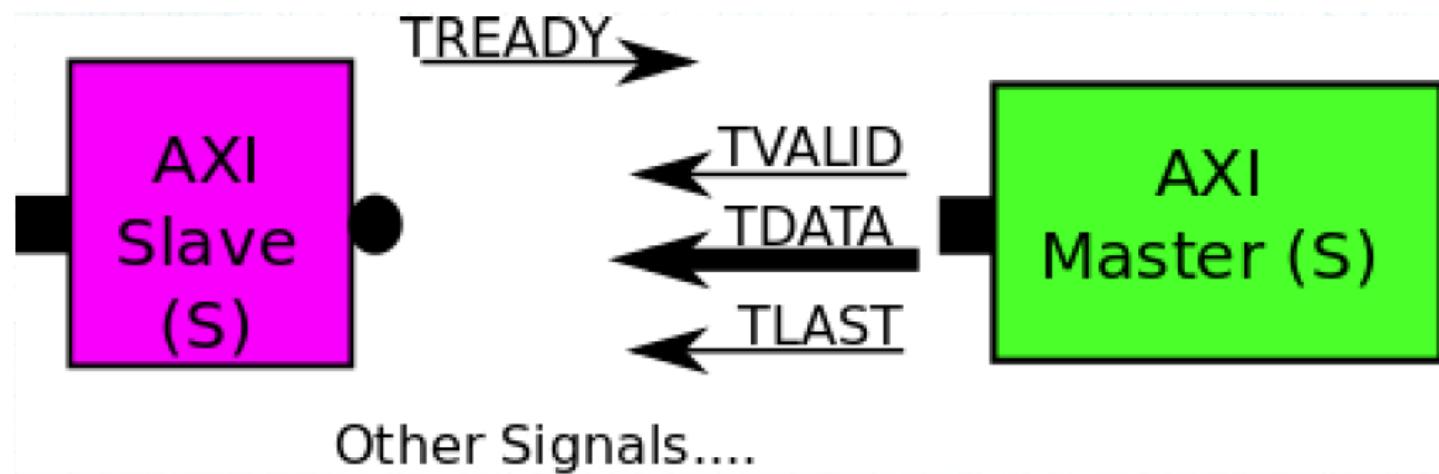
Source: M.S. Sadri, Zynq Training

AXI Memory-Mapped vs. AXI Stream



Source: M.S. Sadri, Zynq Training

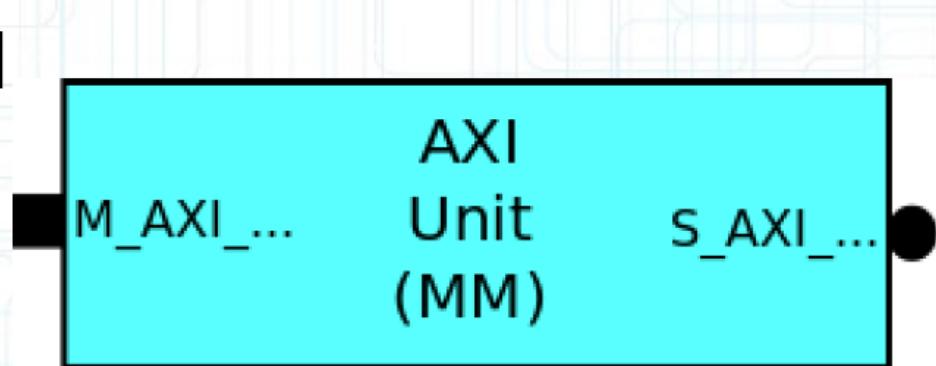
Selected AXI Stream Ports



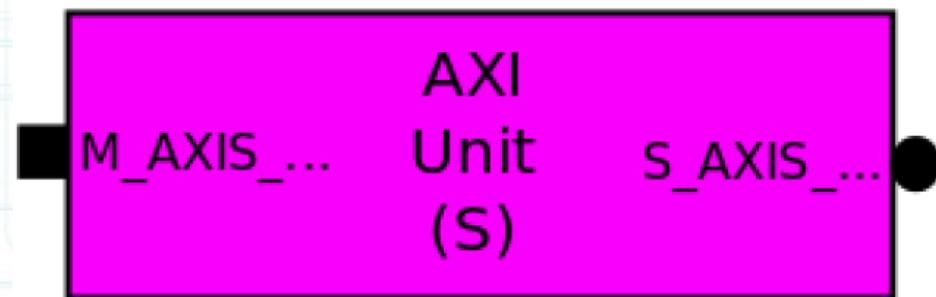
Source: M.S. Sadri, Zynq Training

AXI Port Naming Conventions

- Memory Mapped



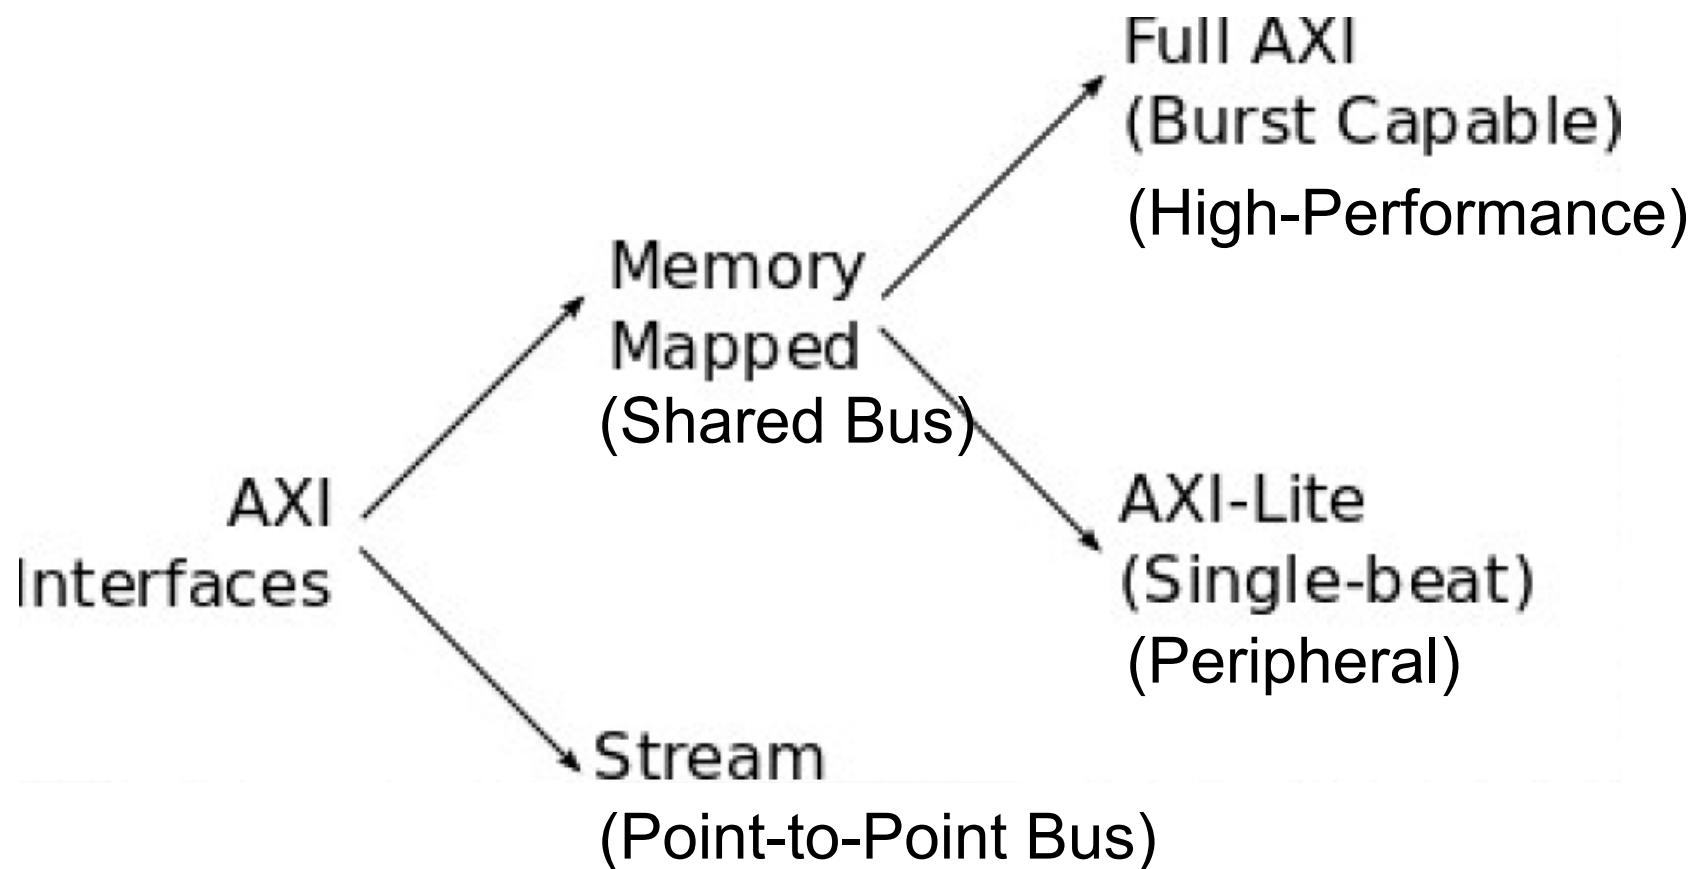
- Stream



- Signal names:

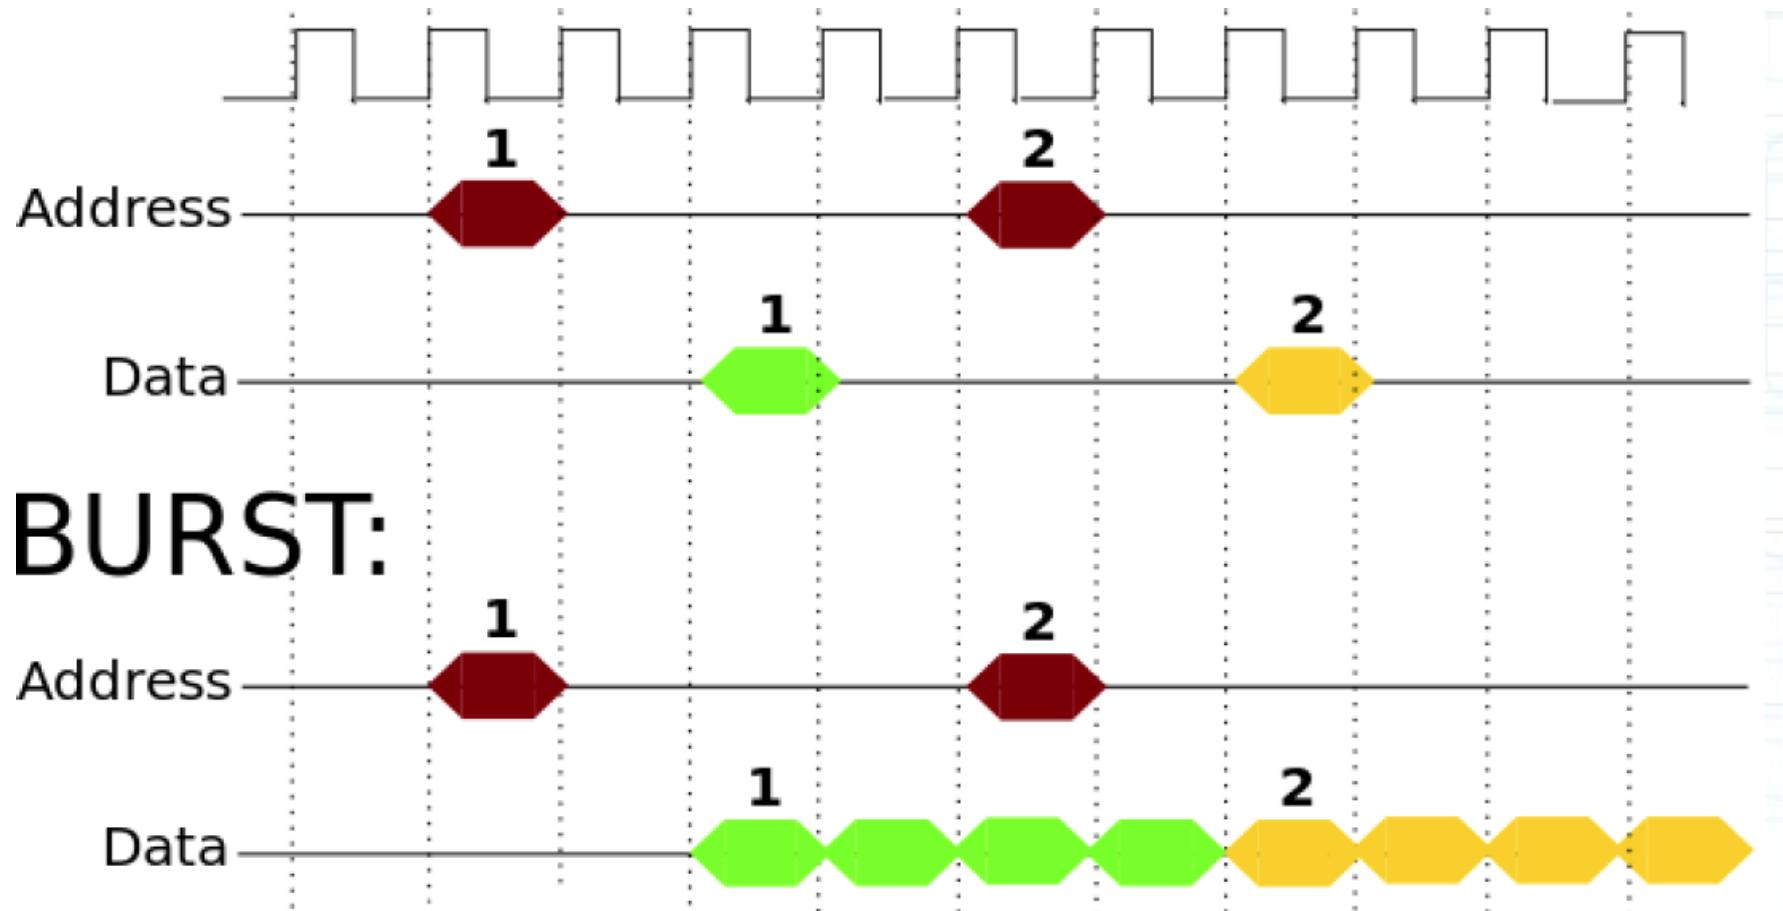
- e.g. S_AXI_awid, M_AXI_arready

AXI Interfaces



Source: M.S. Sadri, Zynq Training

Concept of a Burst



Source: M.S. Sadri, Zynq Training

Competing System-on-Chip Bus Standards

Bus	Developed by	High-Performance Shared Bus	Peripheral Shared Bus	Point-to-Point Bus
AMBA v3	ARM	AHB	APB	
AMBA v4	ARM	AXI4	AXI4-Lite	AXI4-Stream
Coreconnect	IBM	PLB	OPB	
Wishbone	SiliCore Corp.	Crossbar Topology	Shared Topology	Point to Point Topology
Avalon	Altera	Avalon-MM	Avalon-MM	Avalon-ST

AMBA: Advanced Microcontroller Bus Architecture

AXI: Advanced eXtensible Interface

AHB: AMBA High-speed Bus

APB: AMBA Peripheral Bus

PLB: Processor Local Bus

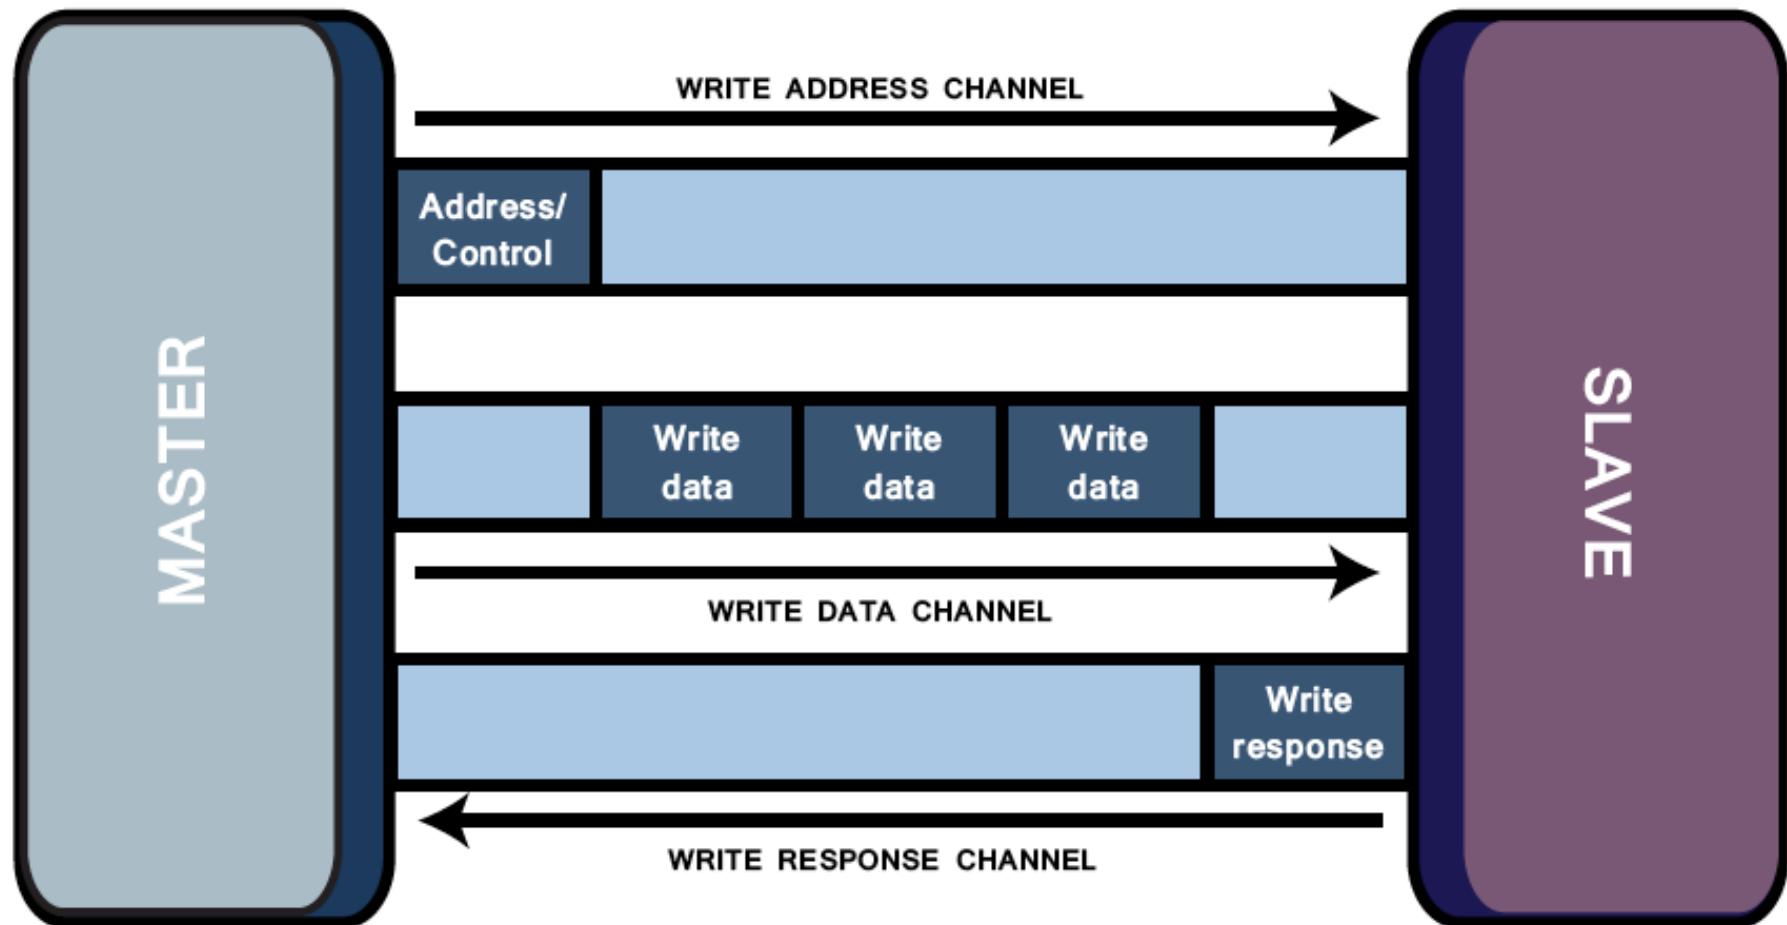
OPB: On-chip Peripheral Bus

MM: Memory Mapped

ST: Streaming

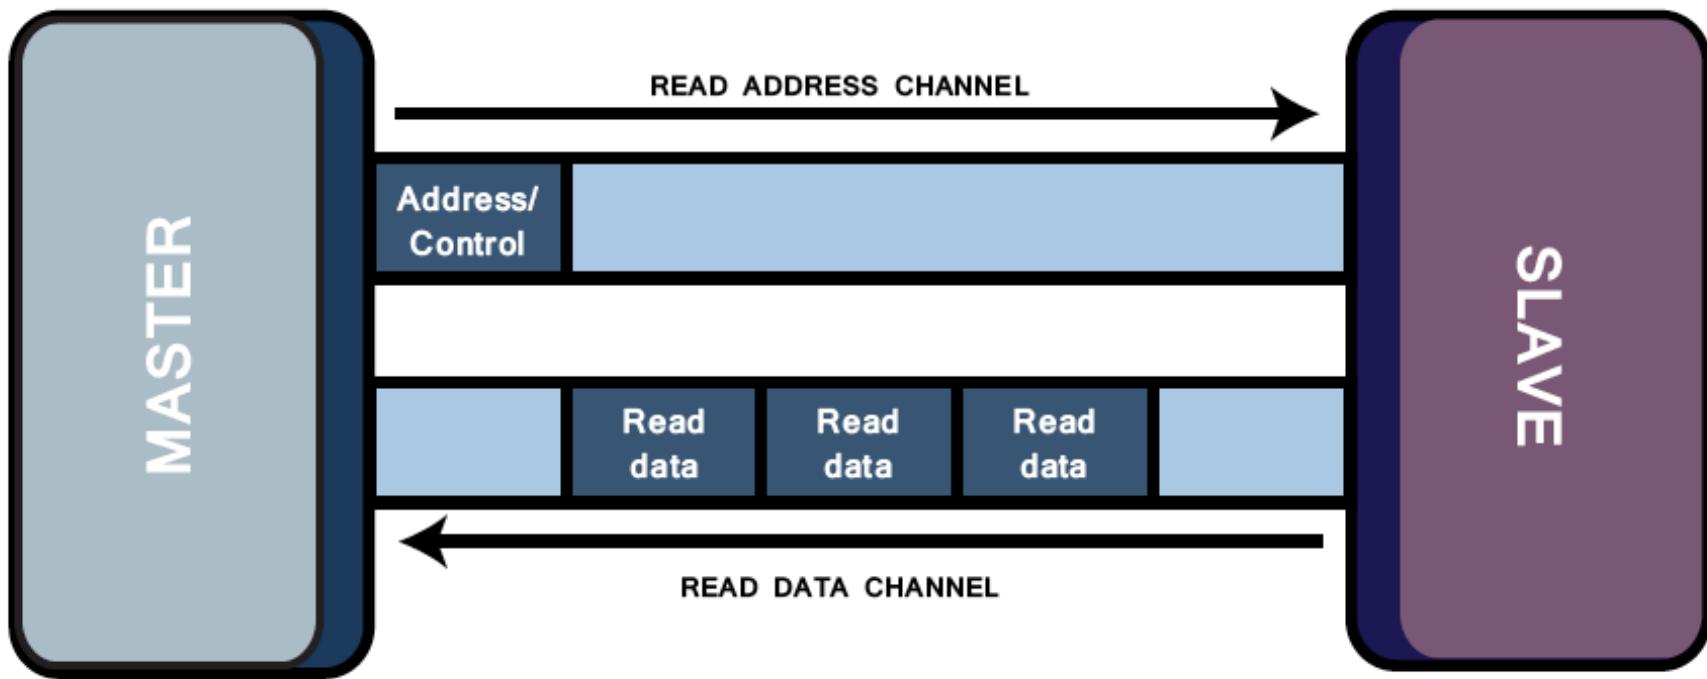
Source: A Practical Introduction to Hardware/Software Codesign

AXI4 Write



Source: The Zynq Book

AXI4 Read



Source: The Zynq Book

AXI4 Interface

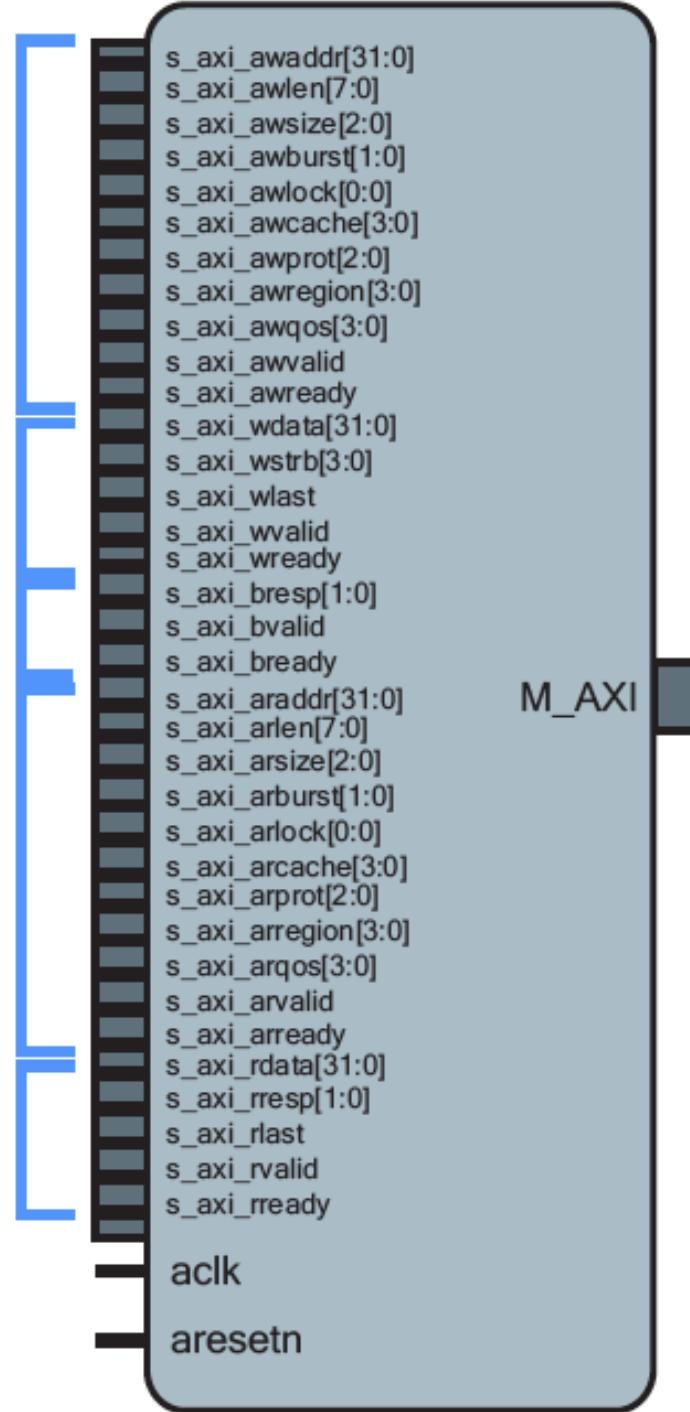
Write Address Channel

Write Data Channel

Write Response Channel

Read Address Channel

Read Data Channel

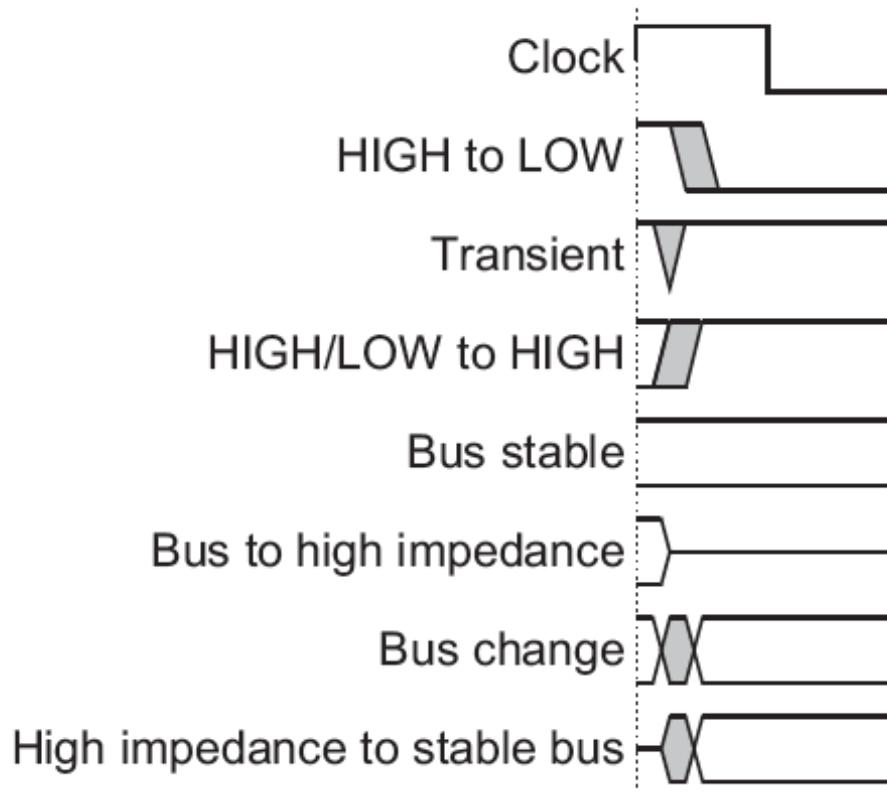


Source: The Zynq Book

Prefixes of Ports from Particular Channels

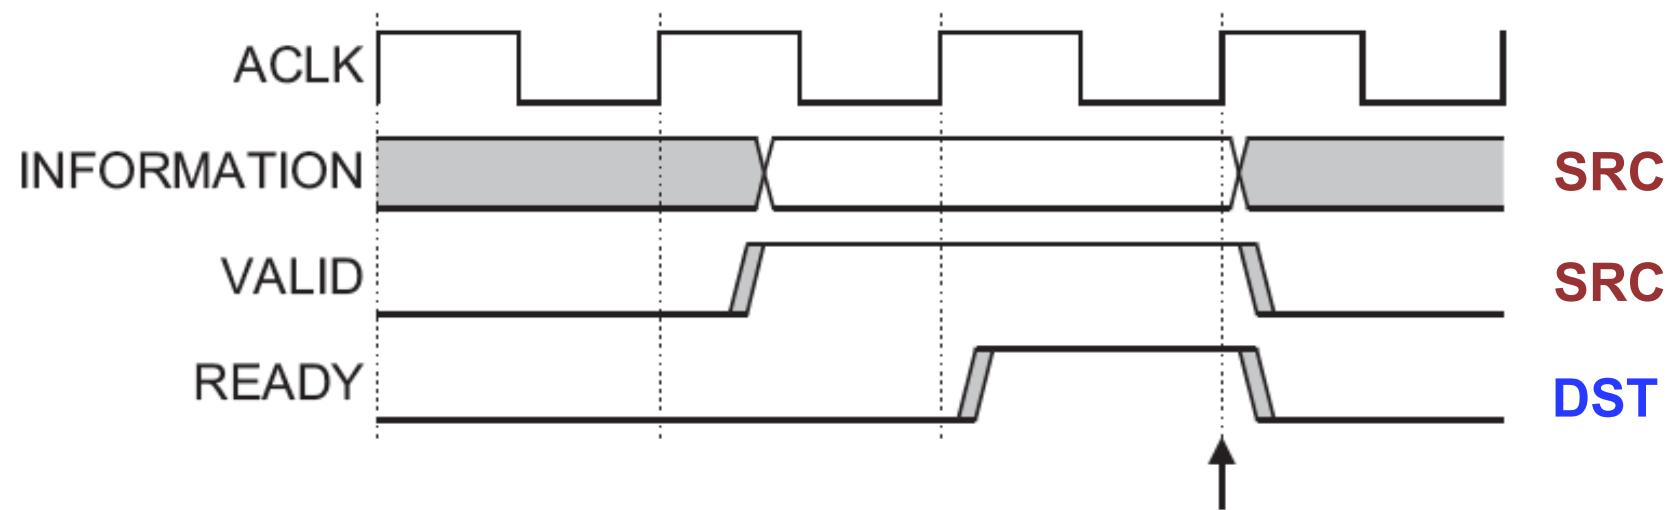
- **Write Address Channel** — the signals contained within this channel are named in the format *s_axi_aw...*
- **Write Data Channel** — the signals contained within this channel are named in the format *s_axi_w...*
- **Write Response Channel** — the signals contained within this channel are named in the format *s_axi_b...*
- **Read Address Channel** — the signals contained within this channel are named in the format *s_axi_ar...*
- **Read Data Channel** — the signals contained within this channel are named in the format *s_axi_r...*

Timing Diagram Conventions



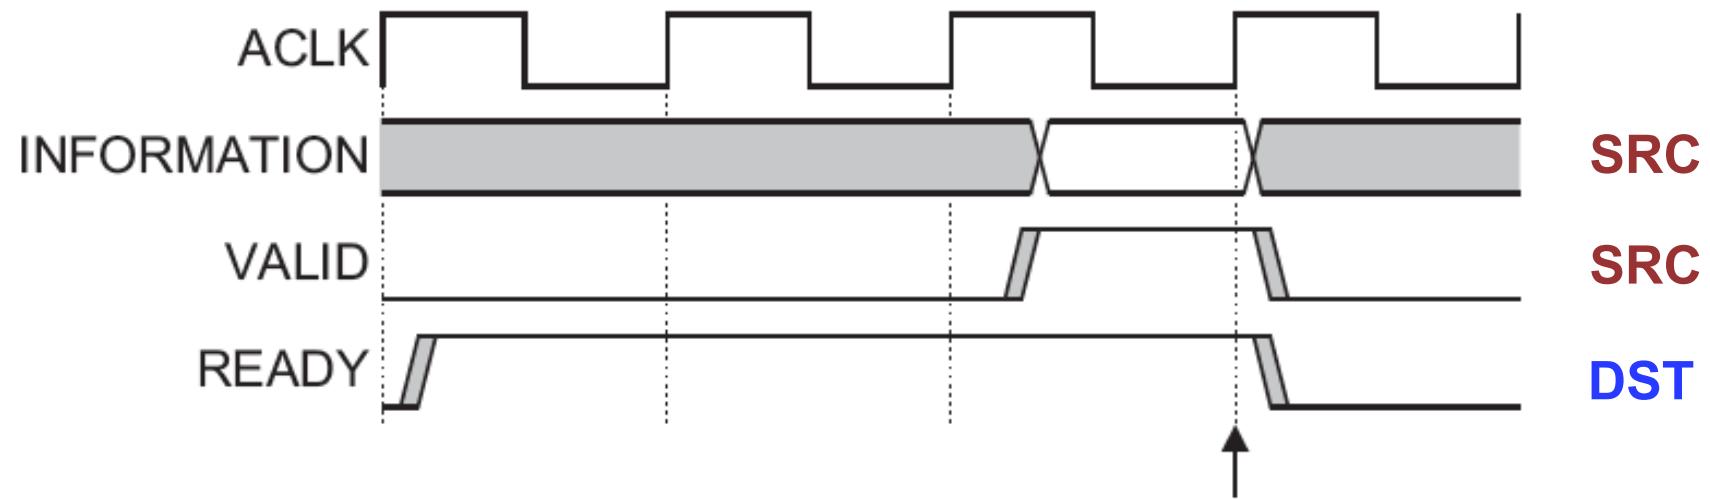
Source: ARM AMBA AXI Protocol v1.0: Specification

VALID before READY Handshake



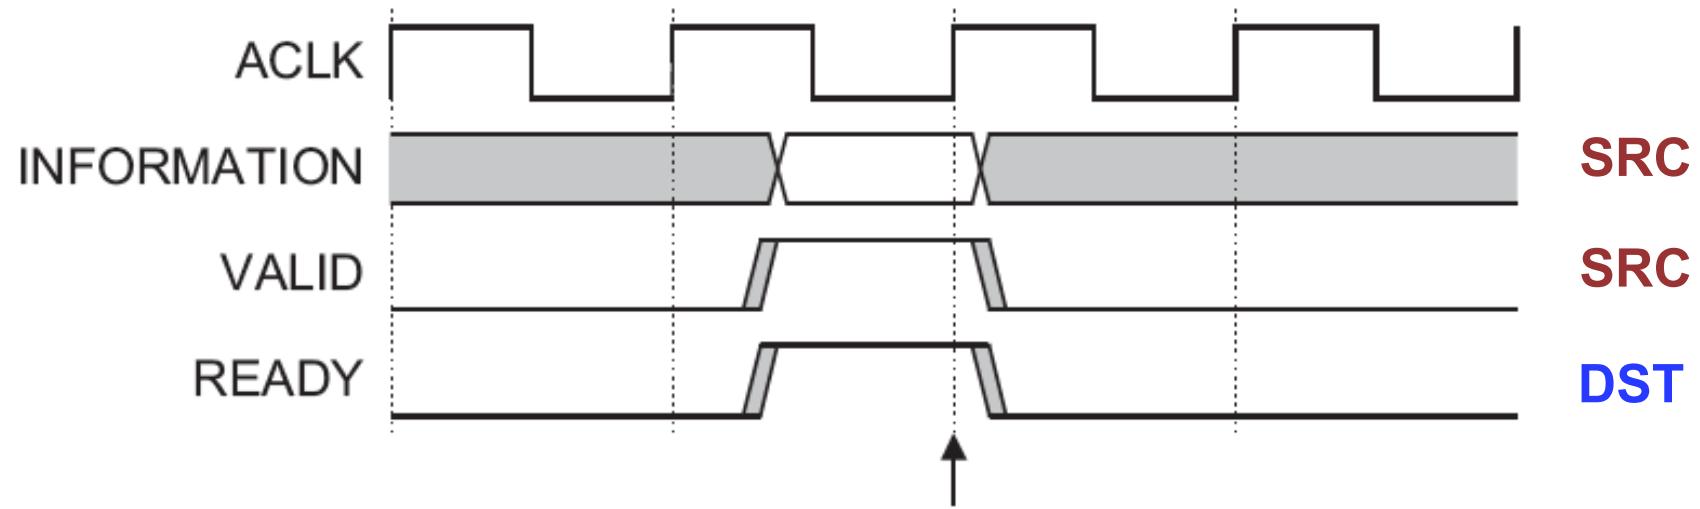
Source: ARM AMBA AXI Protocol v1.0: Specification

READY before VALID Handshake



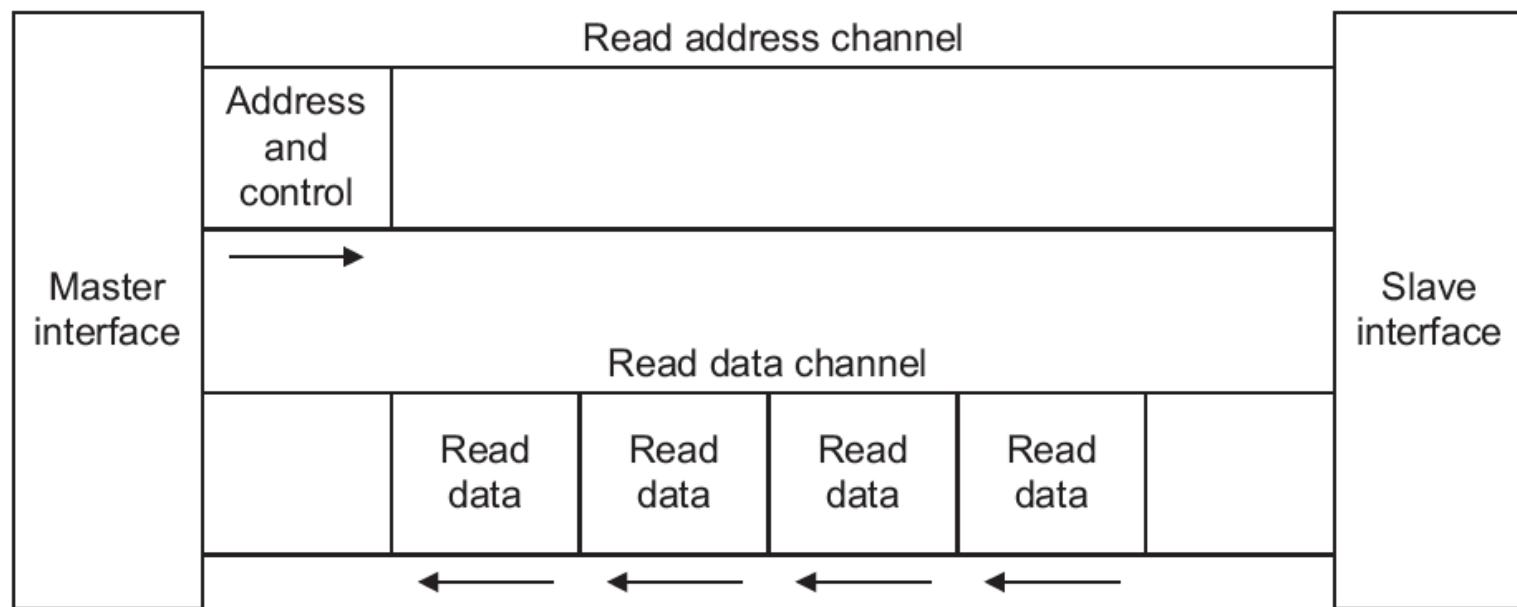
Source: ARM AMBA AXI Protocol v1.0: Specification

VALID with READY Handshake



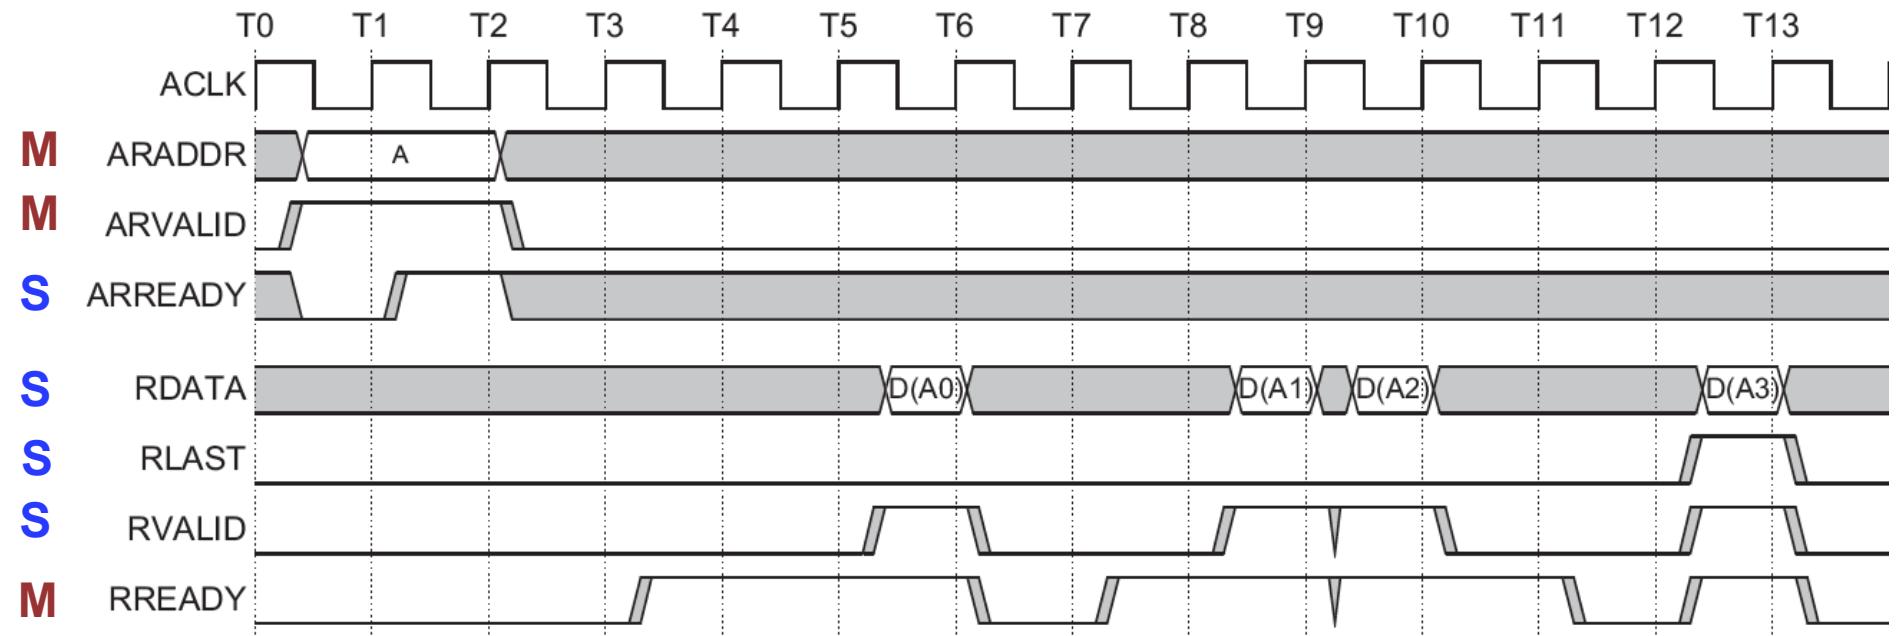
Source: ARM AMBA AXI Protocol v1.0: Specification

Channel Architecture of Reads



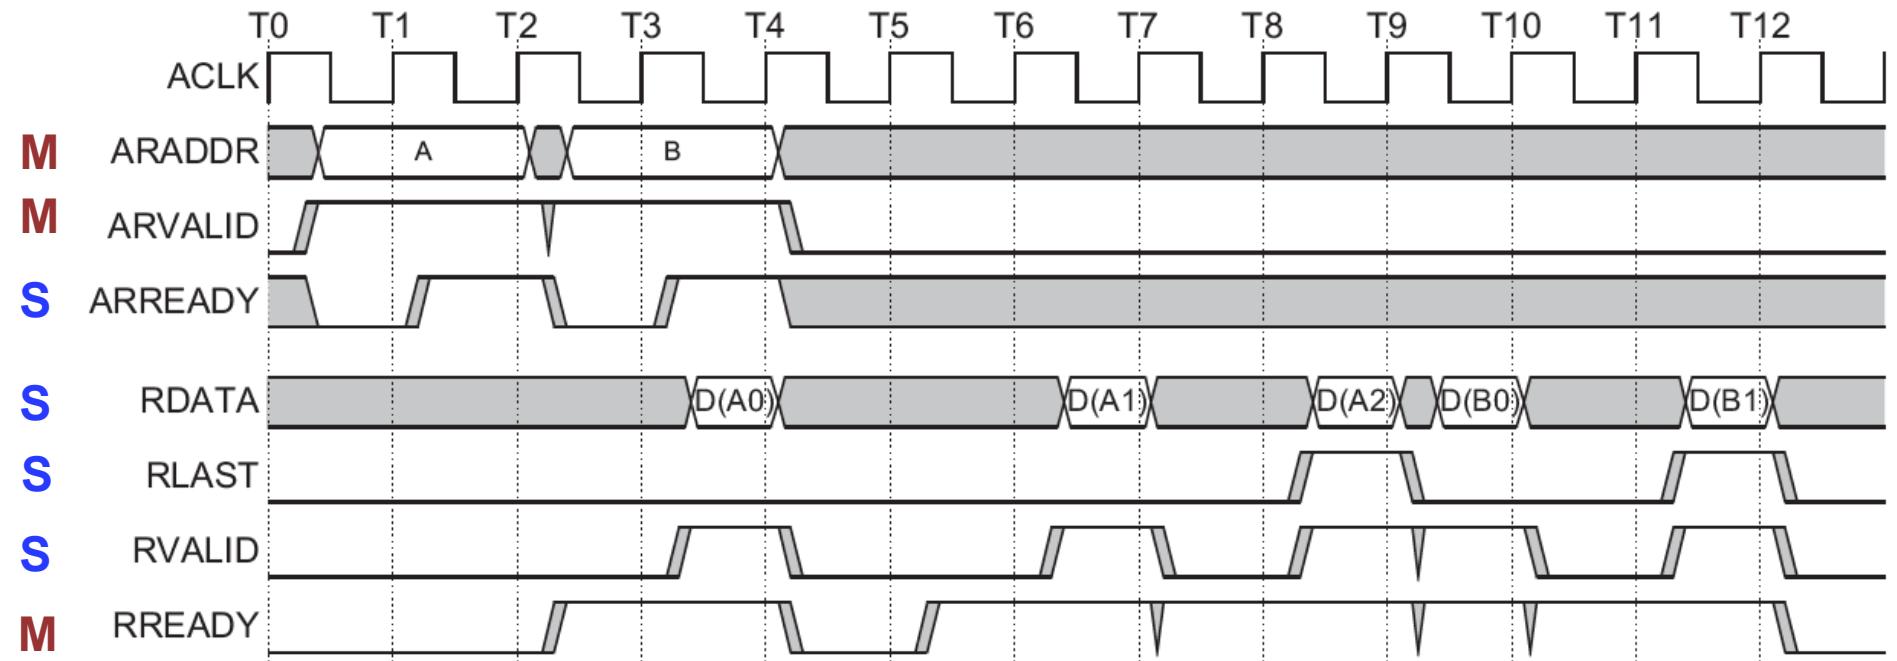
Source: ARM AMBA AXI Protocol v1.0: Specification

Read Burst



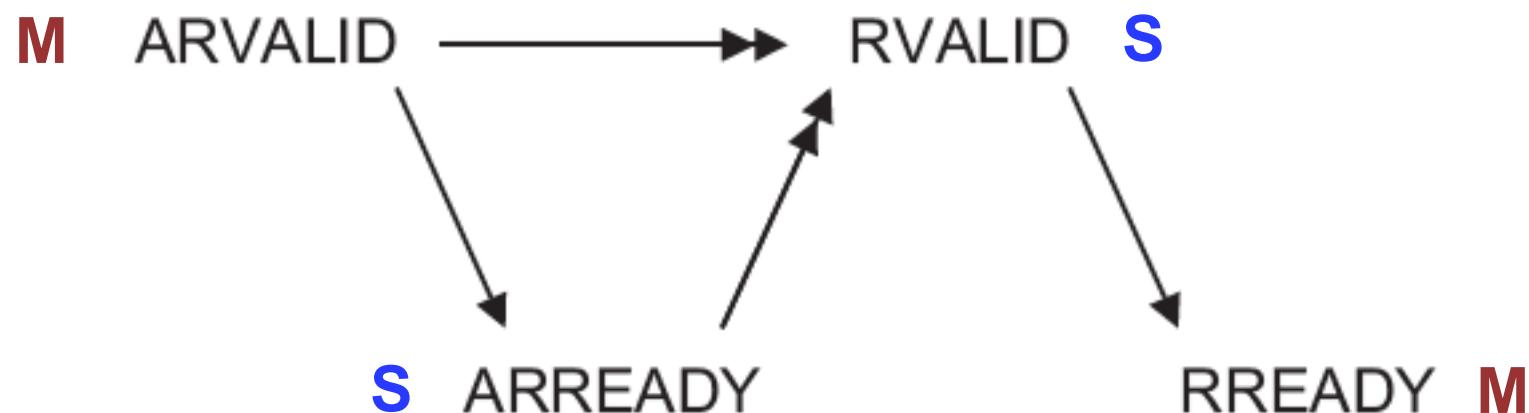
Source: ARM AMBA AXI Protocol v1.0: Specification

Overlapping Read Bursts

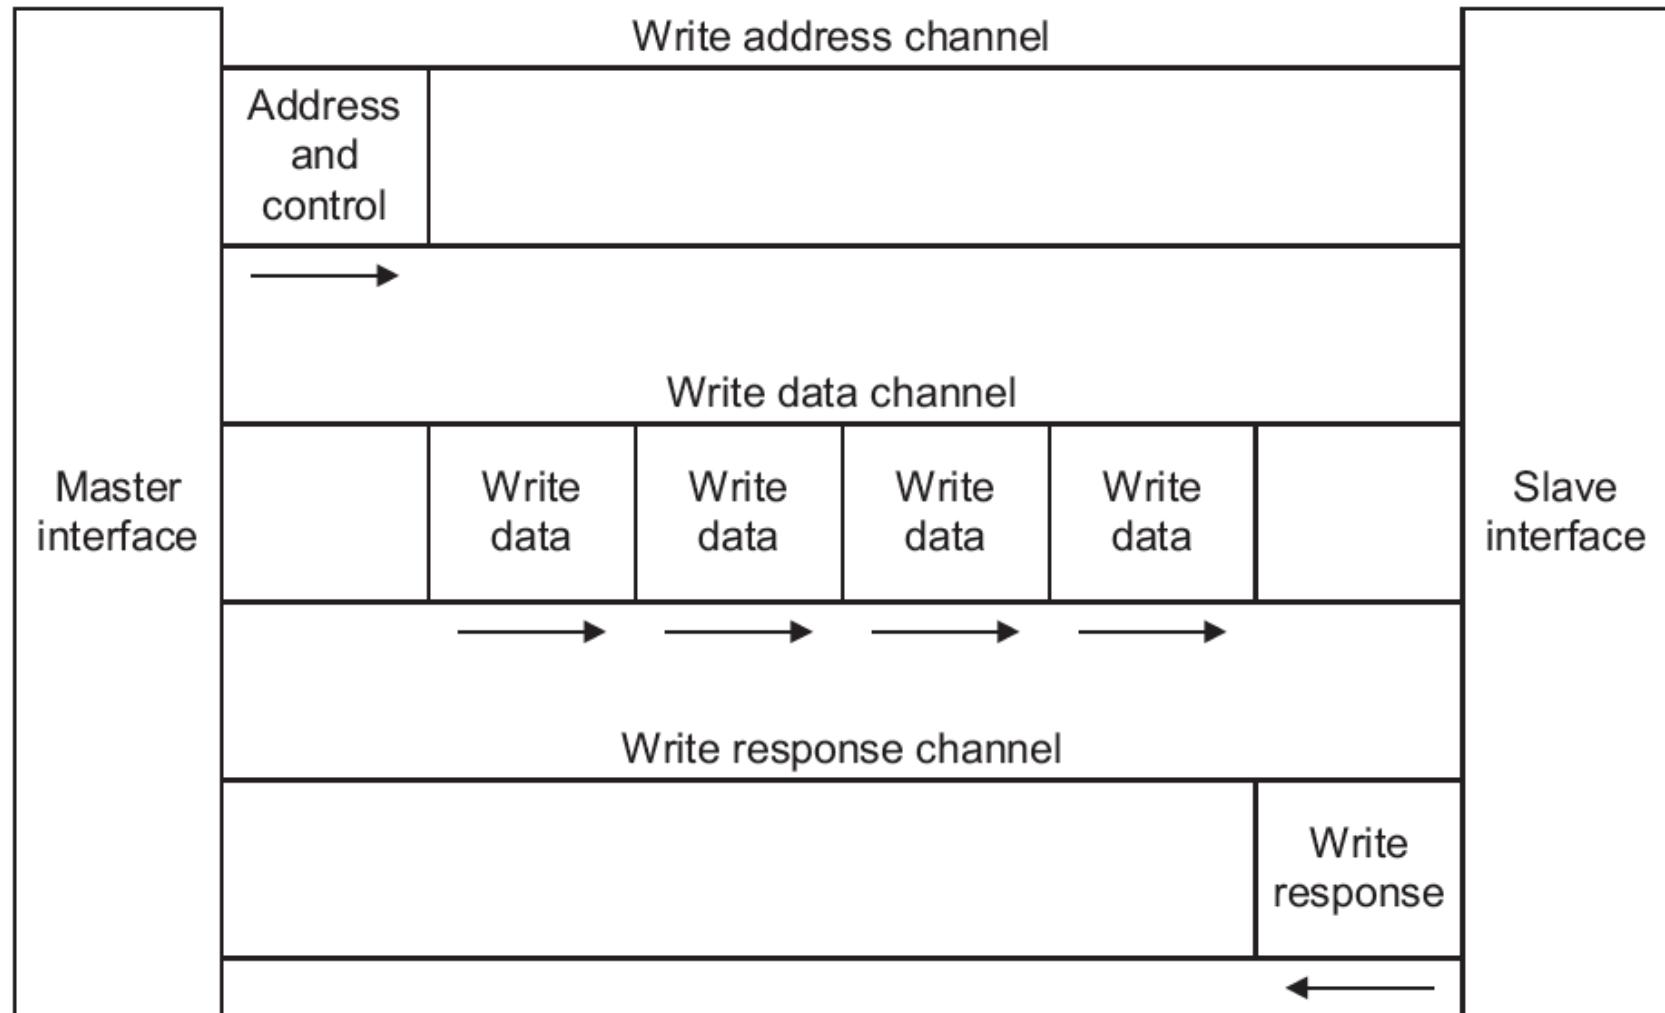


Source: ARM AMBA AXI Protocol v1.0: Specification

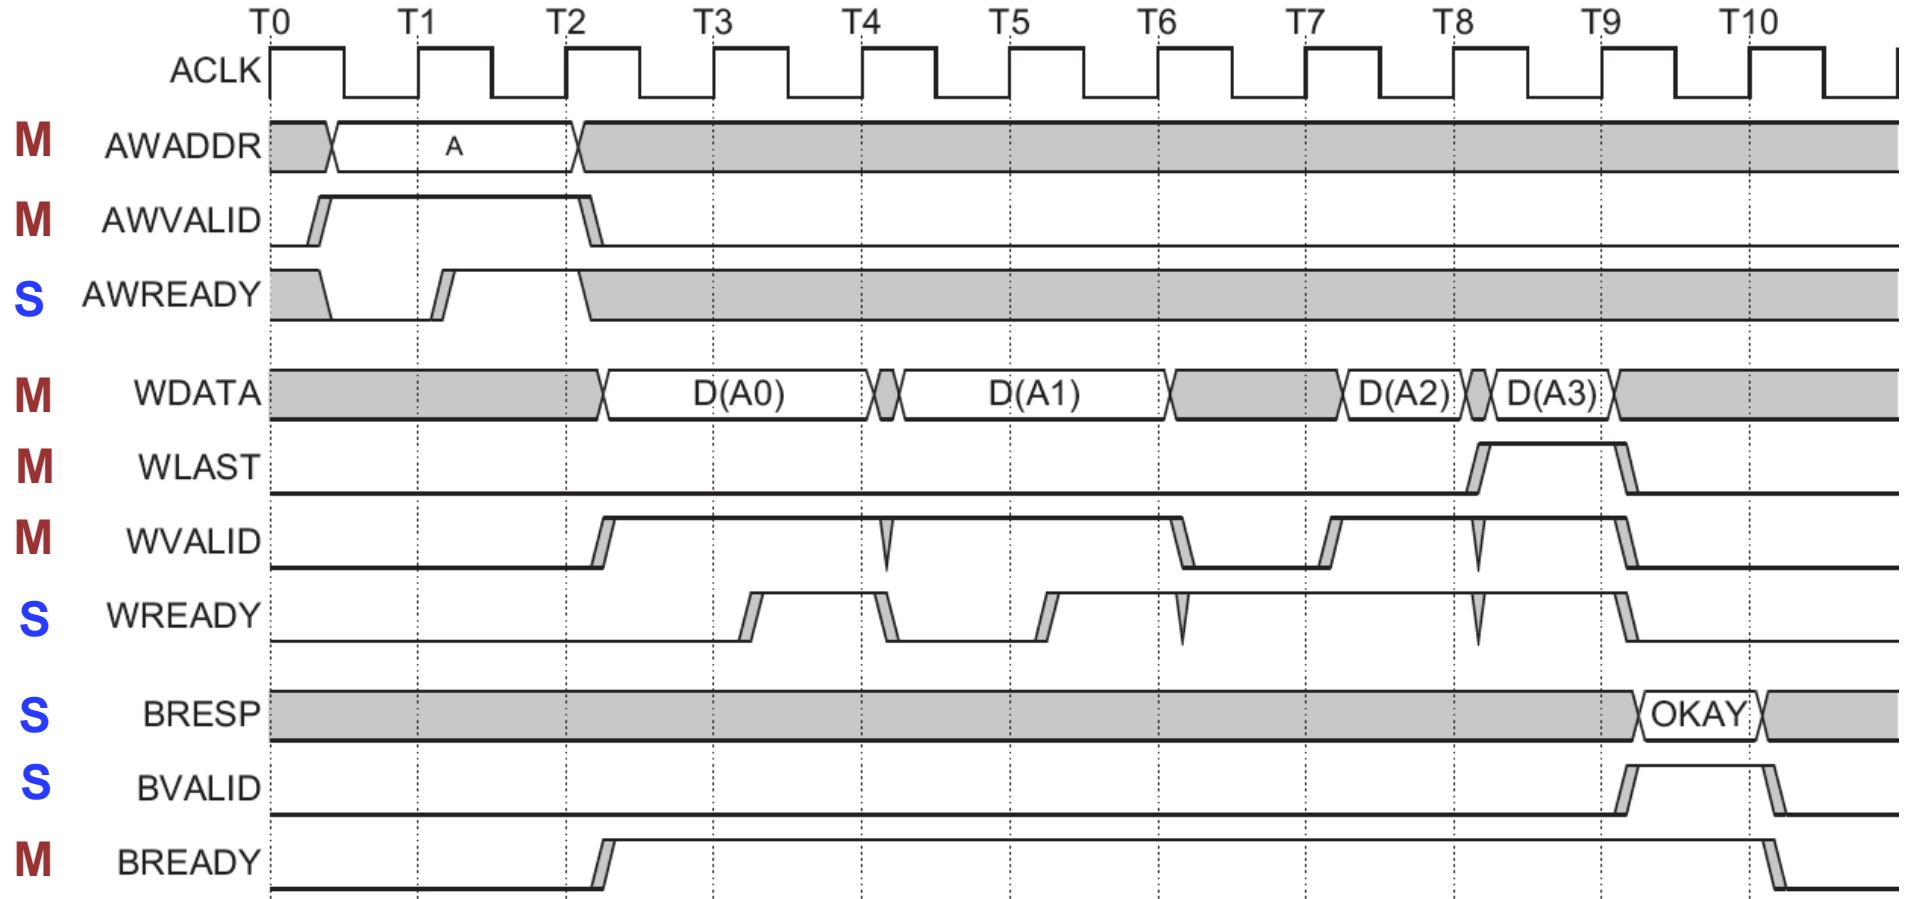
Read Transaction Handshake Dependencies



Channel Architecture of Writes

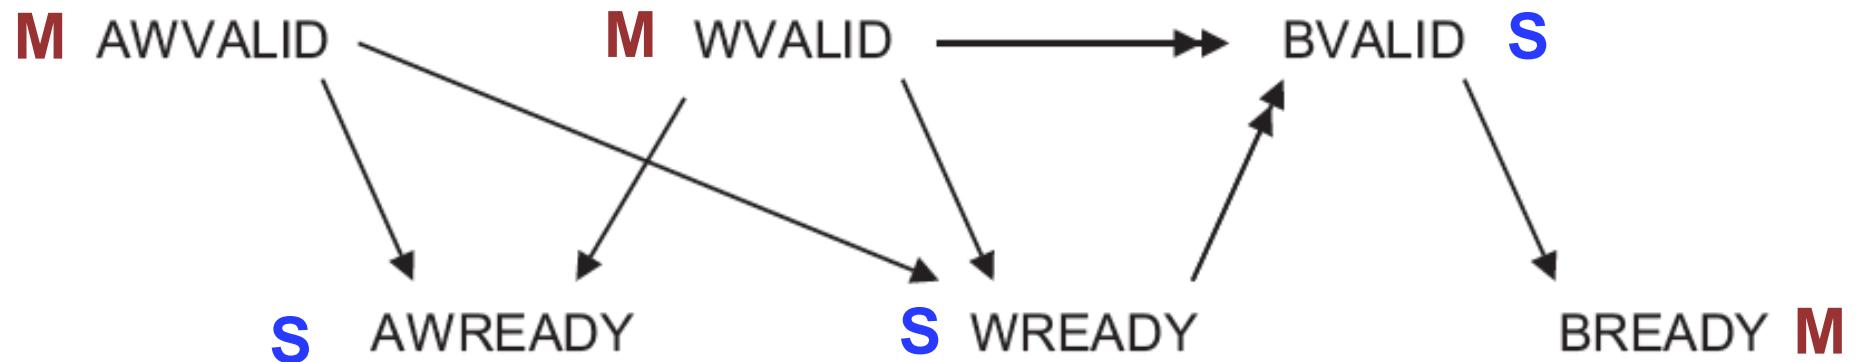


Write Burst



Source: ARM AMBA AXI Protocol v1.0: Specification

Write Transaction Handshake Dependencies

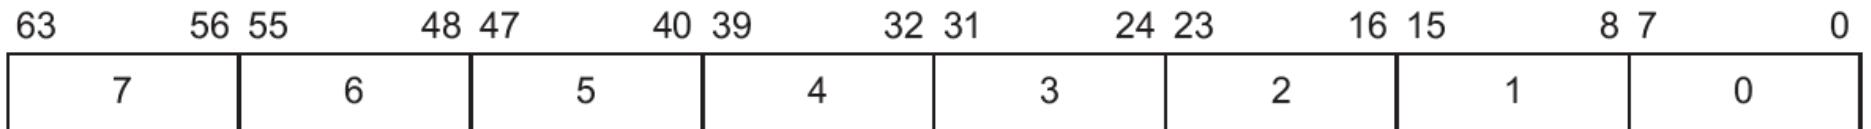


Source: ARM AMBA AXI Protocol v1.0: Specification

ARBURST[1:0]	Burst type	Description	Access
b00	FIXED	Fixed-address burst	FIFO-type
b01	INCR	Incrementing-address burst	Normal sequential memory
b10	WRAP	Incrementing-address burst that wraps to a lower address at the wrap boundary	Cache line
b11	Reserved	-	-

Source: ARM AMBA AXI Protocol v1.0: Specification

Role of Write Strobe WSTRB



WSTRB[n] corresponds to
WDATA[8*n+7 downto 8*n]

Narrow Transfer Example with 8-bit Transfers

Byte lane used

			DATA[7:0]	1st transfer
		DATA[15:8]		2nd transfer
	DATA[23:16]			3rd transfer
DATA[31:24]				4th transfer
			DATA[7:0]	5th transfer

Source: ARM AMBA AXI Protocol v1.0: Specification

Narrow Transfer Example with 32-bit Transfers

Byte lane used	
DATA[63:32]	1st transfer
	DATA[31:0]
DATA[63:32]	3rd transfer

Source: ARM AMBA AXI Protocol v1.0: Specification

Aligned and Unaligned Word Transfers on a 32-bit Bus

Address: 0x00
Transfer size: 32 bits
Burst type: incrementing
Burst length: 4 transfers

31	24 23	16 15	8	7	0	
3		2	1	0		1st transfer
7		6	5	4		2nd transfer
B		A	9	8		3rd transfer
F		E	D	C		4th transfer

Address: 0x01
Transfer size: 32 bits
Burst type: incrementing
Burst length: 4 transfers

3	2	1	0	1st transfer
7	6	5	4	2nd transfer
B	A	9	8	3rd transfer
F	E	D	C	4th transfer

Aligned and Unaligned Word Transfers on a 64-bit Bus

63	56 55	48 47	40 39	32 31	24 23	16 15	8	7	0	
7	6	5	4	3	2	1	0			1st transfer
7	6	5	4	3	2	1	0			2nd transfer
F	E	D	C	B	A	9	8			3rd transfer
F	E	D	C	B	A	9	8			4th transfer
7	6	5	4	3	2	1	0			1st transfer
F	E	D	C	B	A	9	8			2nd transfer
F	E	D	C	B	A	9	8			3rd transfer
17	16	15	14	13	12	11	10			4th transfer

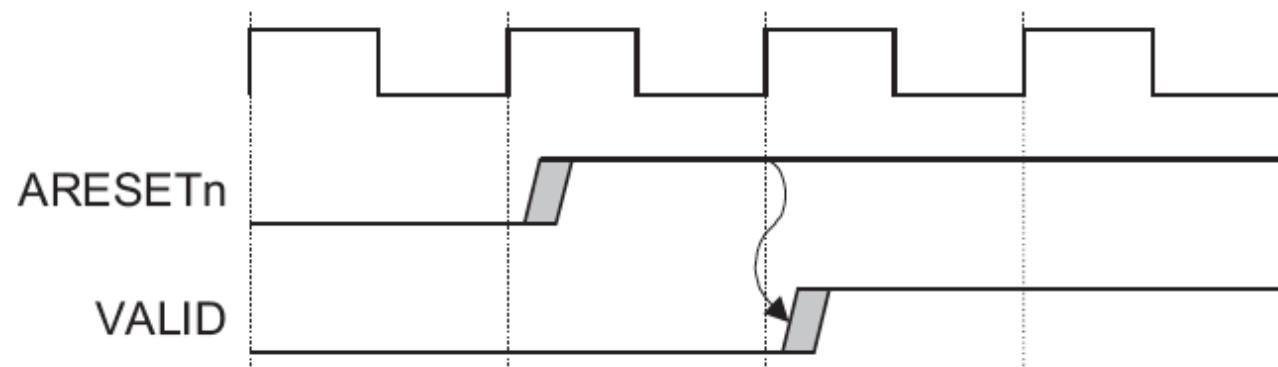
Source: ARM AMBA AXI Protocol v1.0: Specification

Example of IP Core with AXI Interface



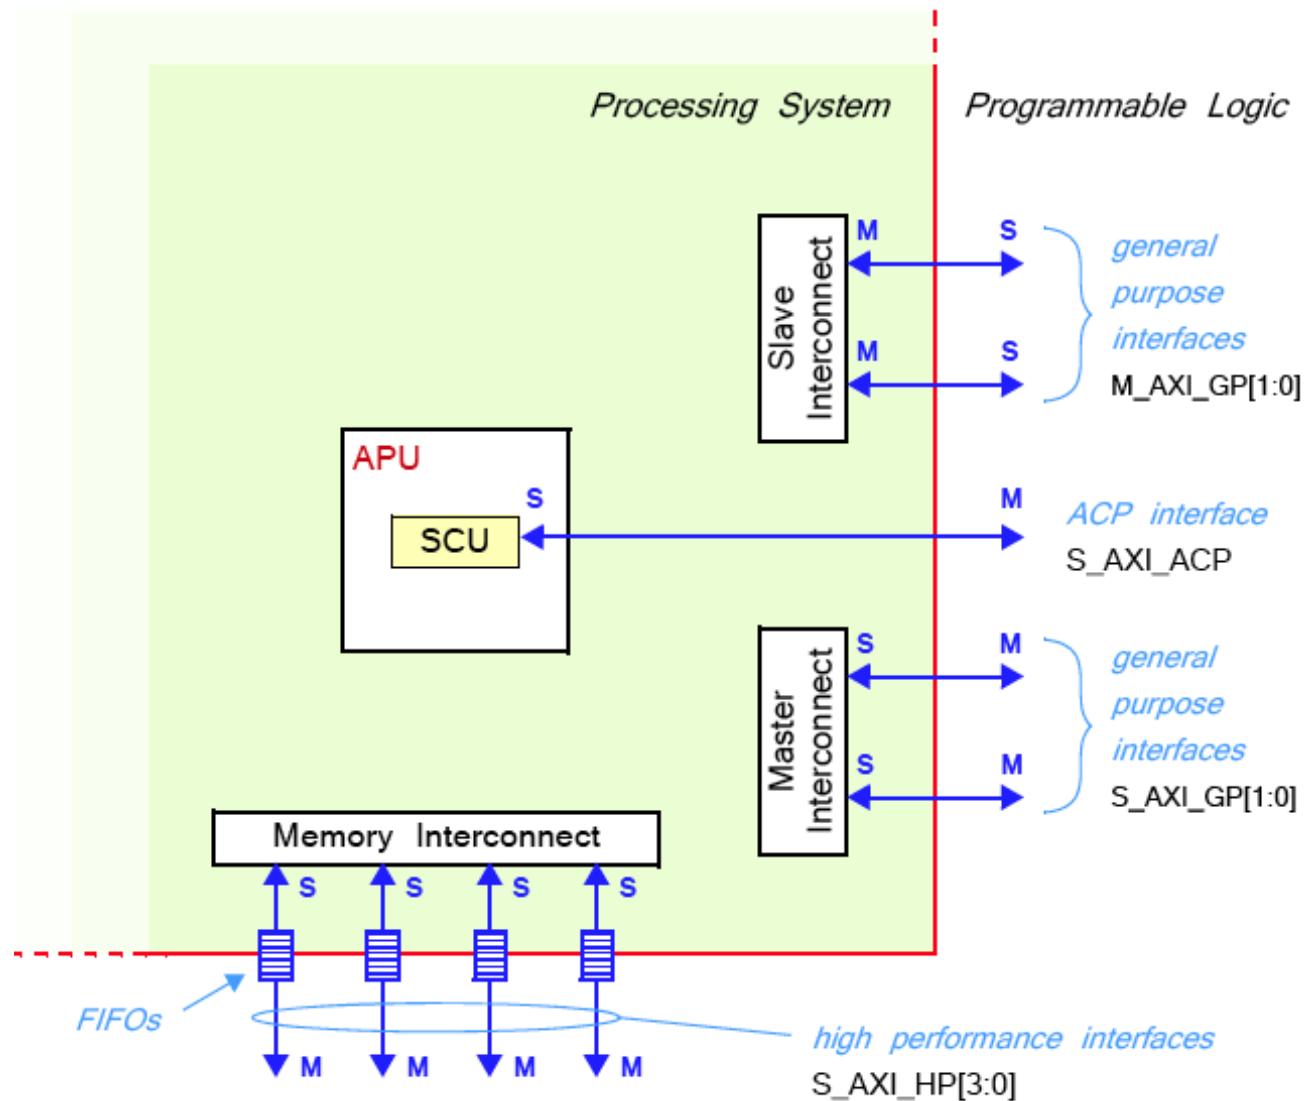
Source: The Zynq Book

Exit from Reset



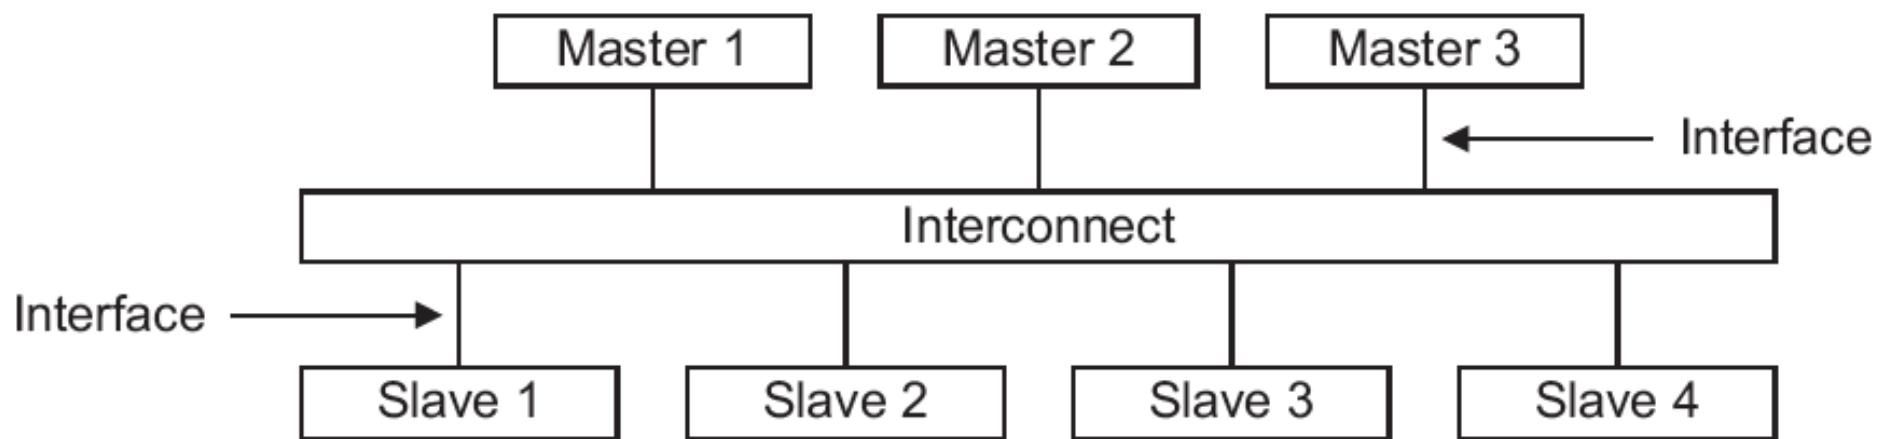
Source: ARM AMBA AXI Protocol v1.0: Specification

PS-PL Interfaces and Interconnects



Source: The Zynq Book

Interconnect vs. Interface



AXI Interfaces and Interconnects

Interface

A point-to-point connection for passing data, addresses, and hand-shaking signals between master and slave clients within the system

Interconnect

A switch which manages and directs traffic between attached AXI interfaces

Zynq AXI PS-PL Interfaces

➤ HP

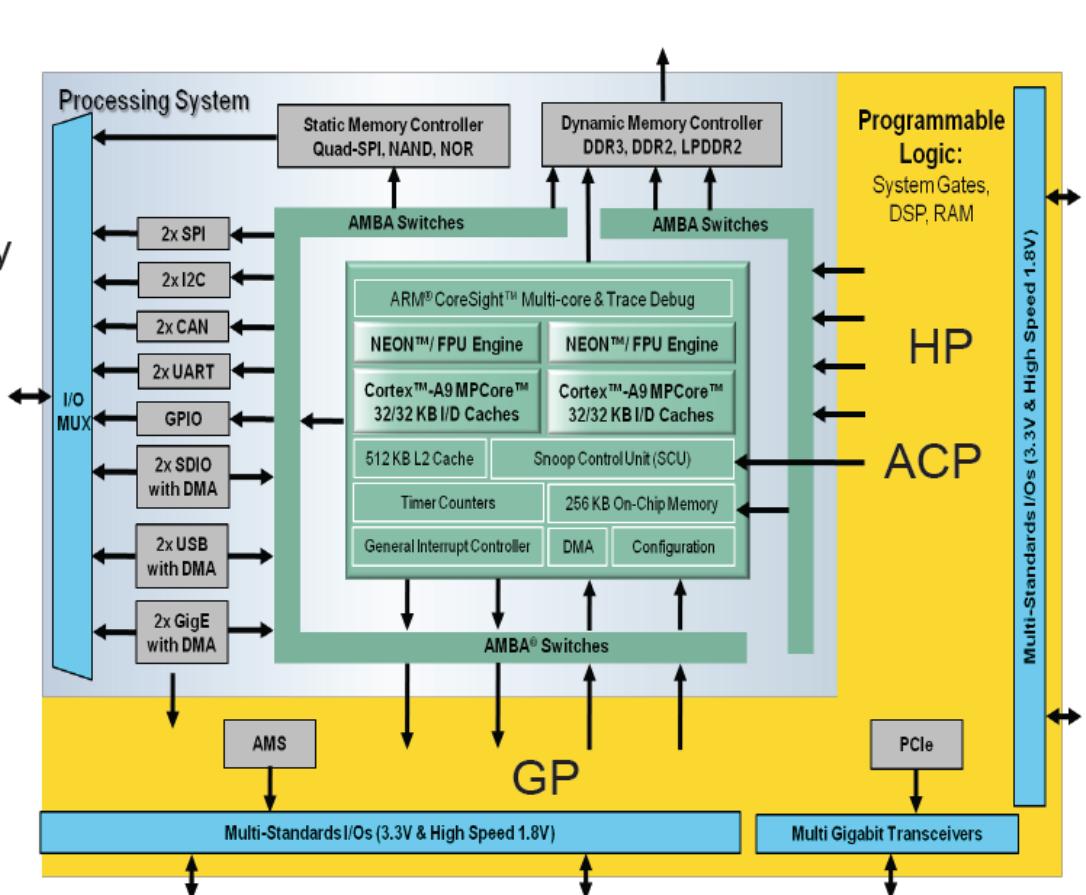
- 4 x 64 bit Slave interfaces
 - Optimized for high bandwidth access from PL to external memory

➤ GP

- 2 x 32 bit Slave interfaces
 - Optimized for access from PL to PS peripherals
- 2 x 32 bit Master interfaces
 - Optimized for access from processors to PL registers

➤ ACP

- 1 x 64 bit Slave interface
 - Optimized for access from PL to processor caches



ACP - Accelerator Coherency Port

AXI Interfaces between PS and PL

Interface Name	Interface Description	Master	Slave
M_AXI_GP0	General Purpose (AXI_GP)	PS	PL
M_AXI_GP1		PS	PL
S_AXI_GP0	General Purpose (AXI_GP)	PL	PS
S_AXI_GP1		PL	PS
S_AXI_ACP	Accelerator Coherency Port (ACP), cache coherent transaction	PL	PS
S_AXI_HP0	(Note that AXI_HP interfaces are sometimes referred to as AXI Fifo Interfaces, or AFIs). High Performance Ports (AXI_HP) with read/write FIFOs.	PL	PS
S_AXI_HP1		PL	PS
S_AXI_HP2		PL	PS
S_AXI_HP3		PL	PS

Source: The Zynq Book

General-Purpose Port Summary

- GP ports are designed for maximum flexibility
- Allow register access from PS to PL or PL to PS
- Good for Synchronization
- Prefer ACP or HP port for data transport

High-Performance Port Summary

- HP ports are designed for maximum bandwidth access to external memory and On-Chip Memory (OCM)
- When combined can saturate external memory and OCM bandwidth
 - HP Ports : $4 * 64 \text{ bits} * 150 \text{ MHz} * 2 = 9.6 \text{ GByte/sec}$
 - external DDR: $1 * 32 \text{ bits} * 1066 \text{ MHz} * 2 = 4.3 \text{ GByte/sec}$
 - OCM : $64 \text{ bits} * 222 \text{ MHz} * 2 = 3.5 \text{ GByte/sec}$
- Optimized for large burst lengths and many outstanding transactions
- Large data buffers to amortize access latency
- Efficient upsizing/downsizing for 32 bit accesses

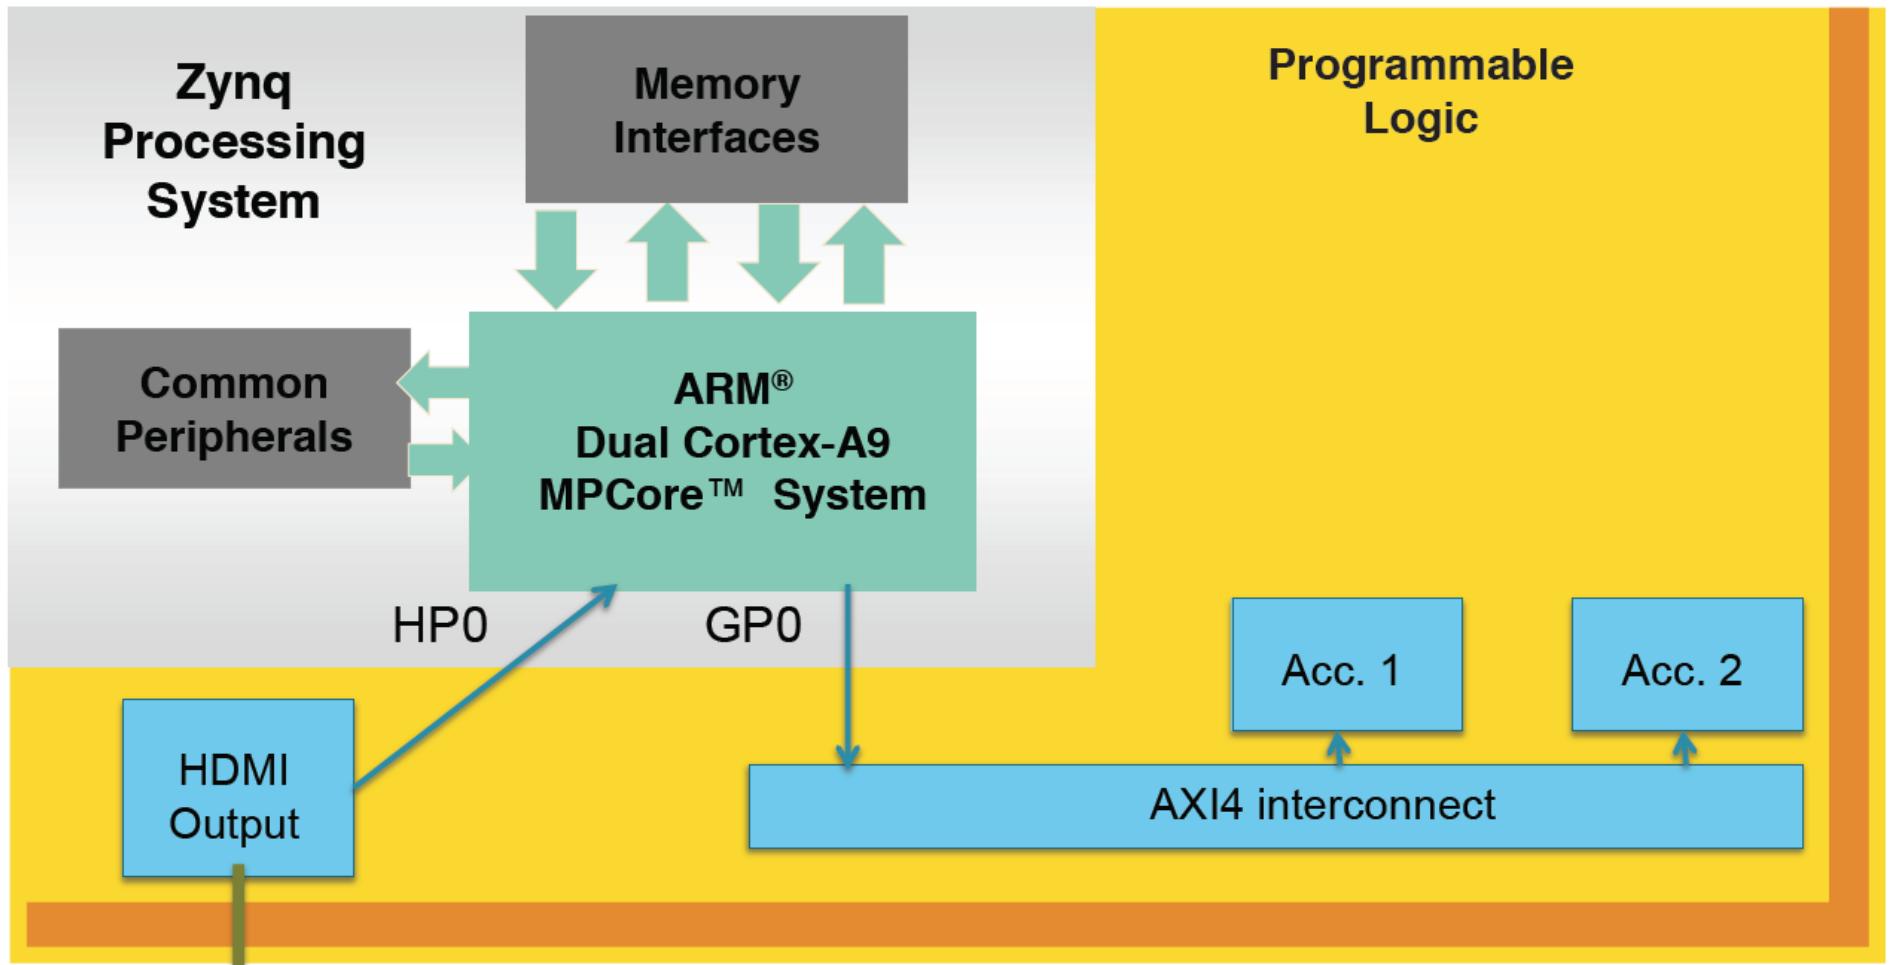
Accelerator Coherency Port (ACP) Summary

- ACP allows limited support for Hardware Coherency
 - Allows a PL accelerator to access cache of the Cortex-A9 processors
 - PL has access through the same path as CPUs including caches, OCM, DDR, and peripherals
 - Access is low latency (assuming data is in processor cache)
no switches in path
- ACP does not allow full coherency
 - PL is not notified of changes in processor caches
 - Use write to PL register for synchronization
- ACP is compromise between bandwidth and latency
 - Optimized for cache line length transfers
 - Low latency for L1/L2 hits
 - Minimal buffering to hide external memory latency
 - One shared 64 bit interface, limit of 8 masters

Accelerator Architecture with Bus Slave

Pro: Simple System Architecture

Con: Limited communication bandwidth

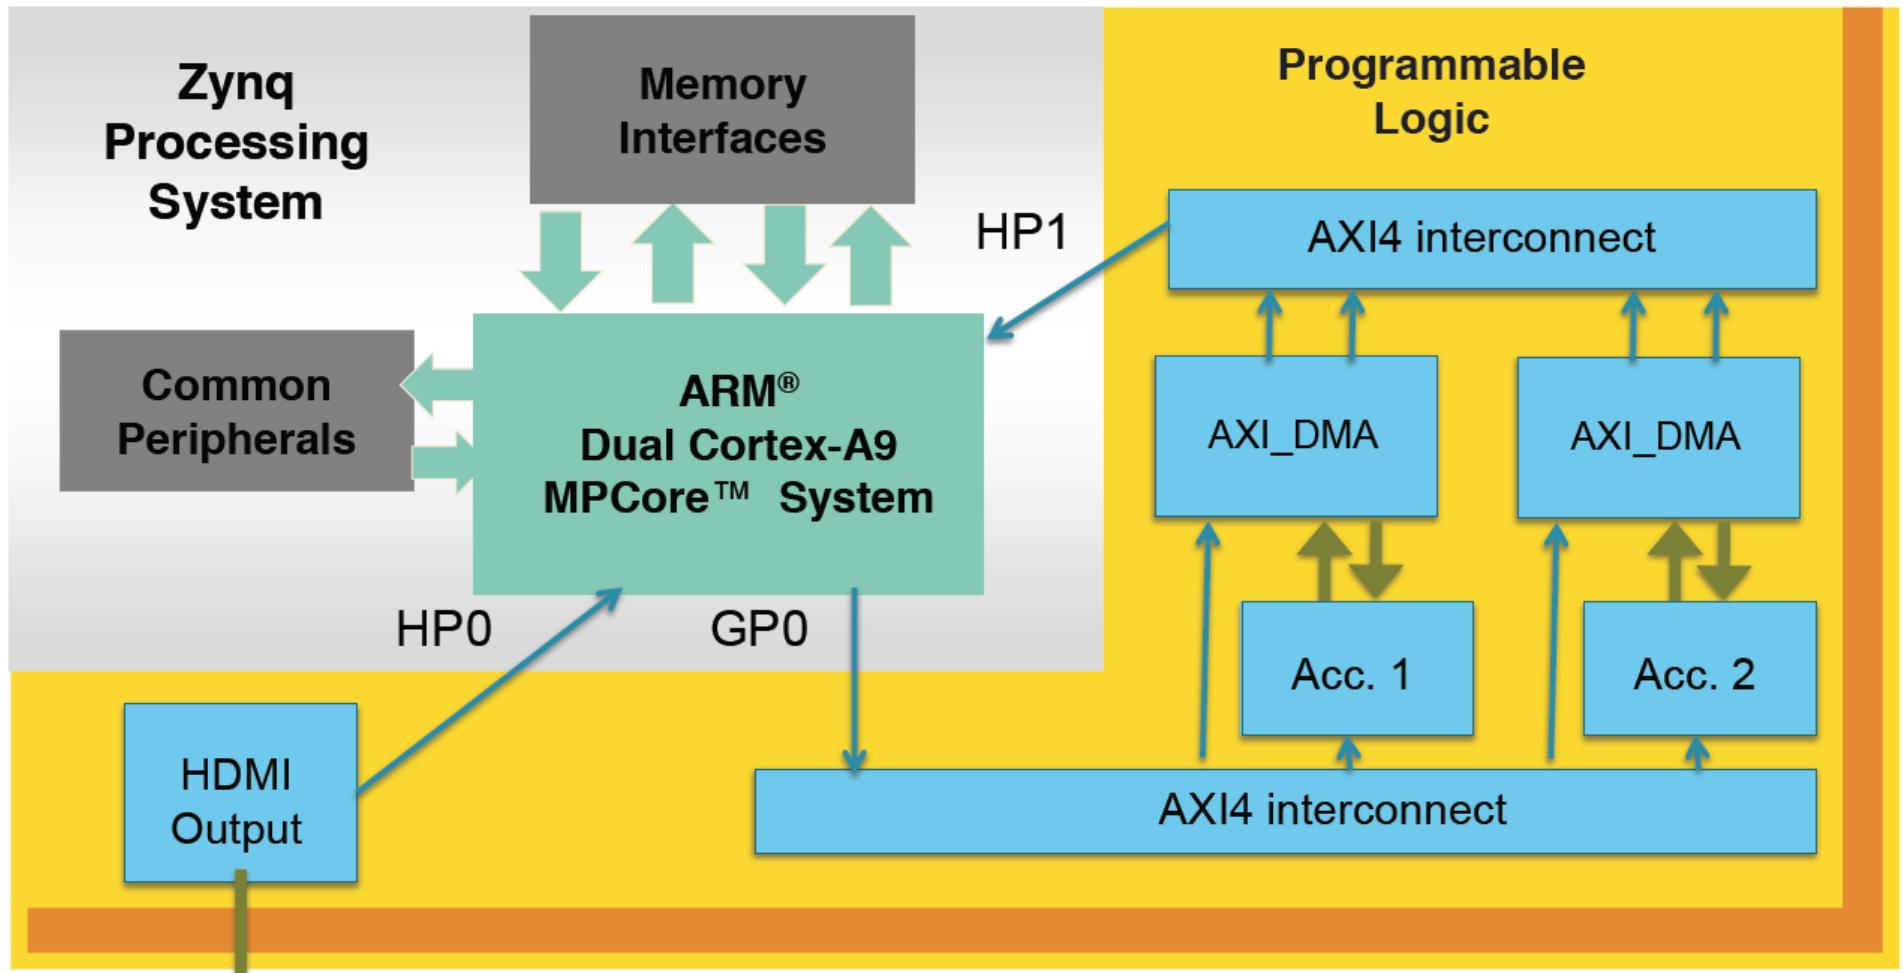


Source: Building Zynq Accelerators with Vivado HLS, FPL 2013 Tutorial

Accelerator Architecture with DMA

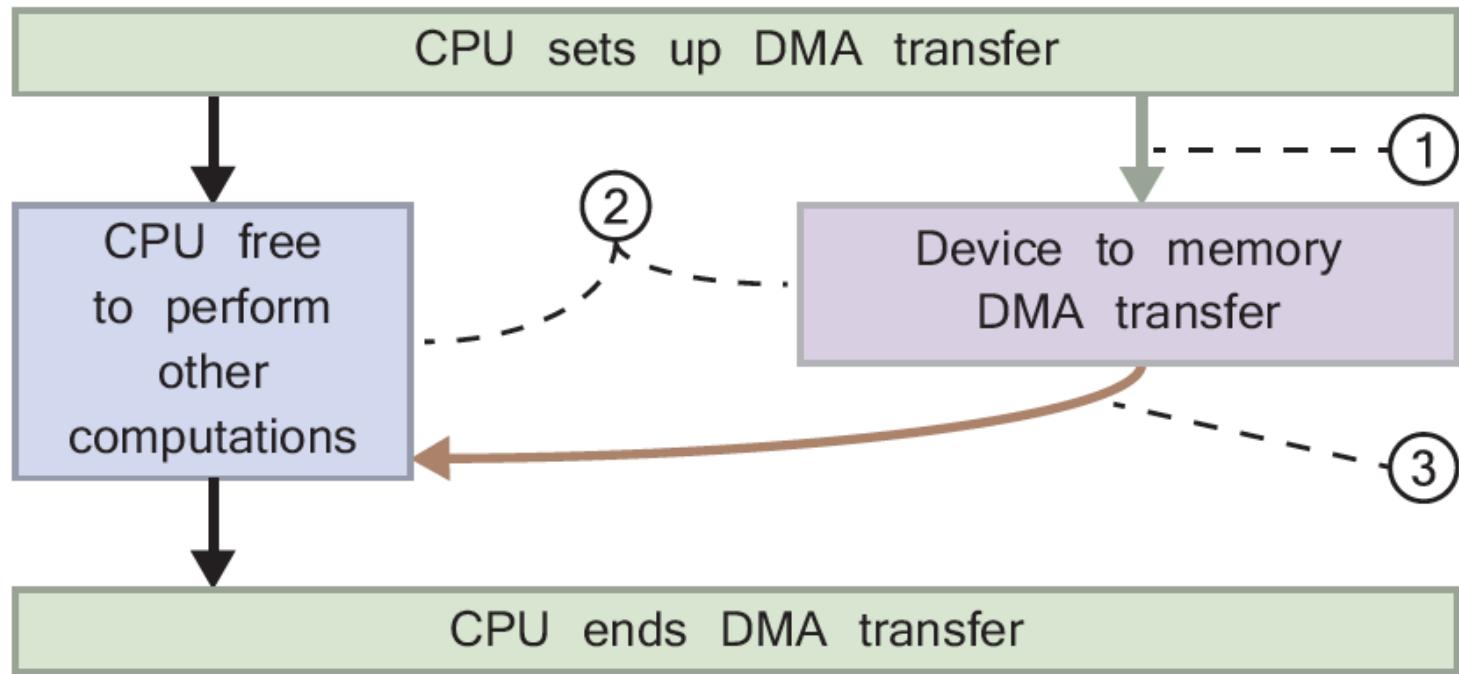
Pro: High Bandwidth Communication

Con: Complicated System Architecture, High Latency



Source: Building Zynq Accelerators with Vivado HLS, FPL 2013 Tutorial

DMA Memory Transfer Operation



Source: The Zynq Book

AXI DMA-based Accelerator Communication

Write to Accelerator

- processor allocates buffer
- processor writes data into buffer
- processor flushes cache for buffer
- processor initiates DMA transfer

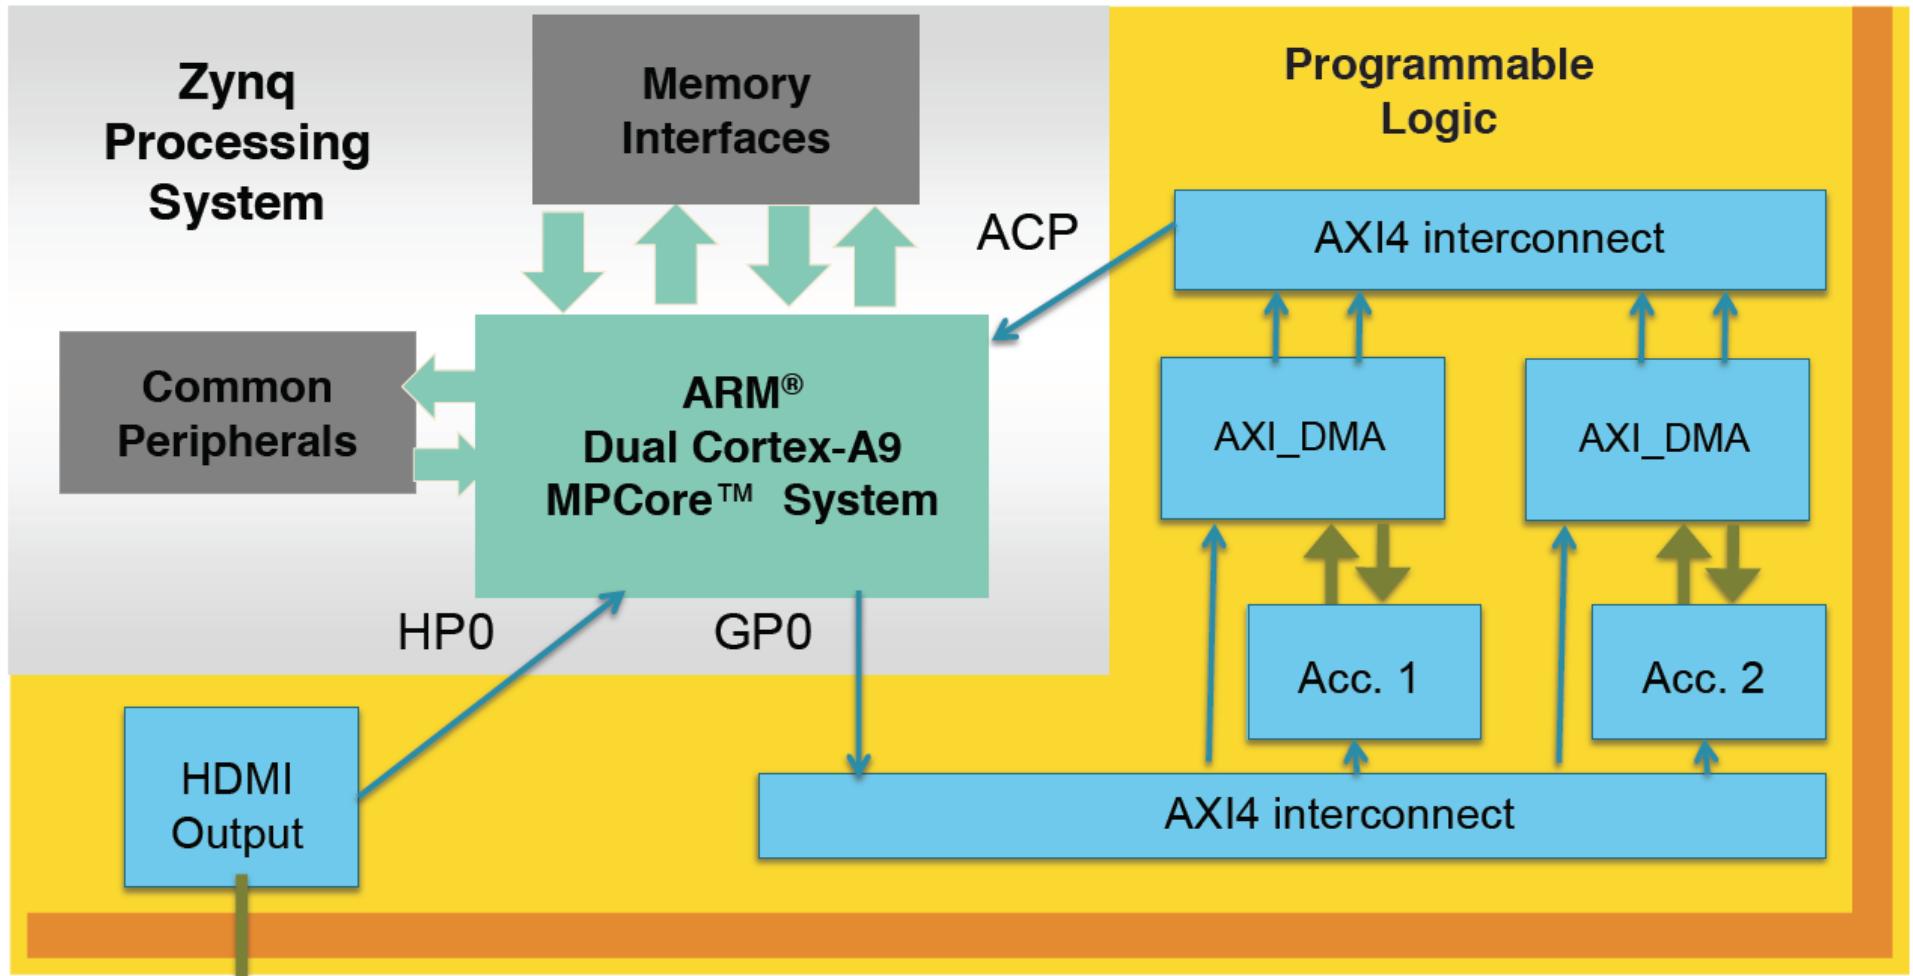
Read from Accelerator

- processor allocates buffer
- processor initiates DMA transfer
- processor waits for DMA to complete
- processor invalidates cache for buffer
- processor reads data from buffer

Accelerator Architecture with Coherent DMA

Pro: Low latency, high-bandwidth communication

Con: Complicated system architecture, Limited to data that fits in caches



Source: Building Zynq Accelerators with Vivado HLS, FPL 2013 Tutorial

Coherent AXI DMA-based Accelerator Communication

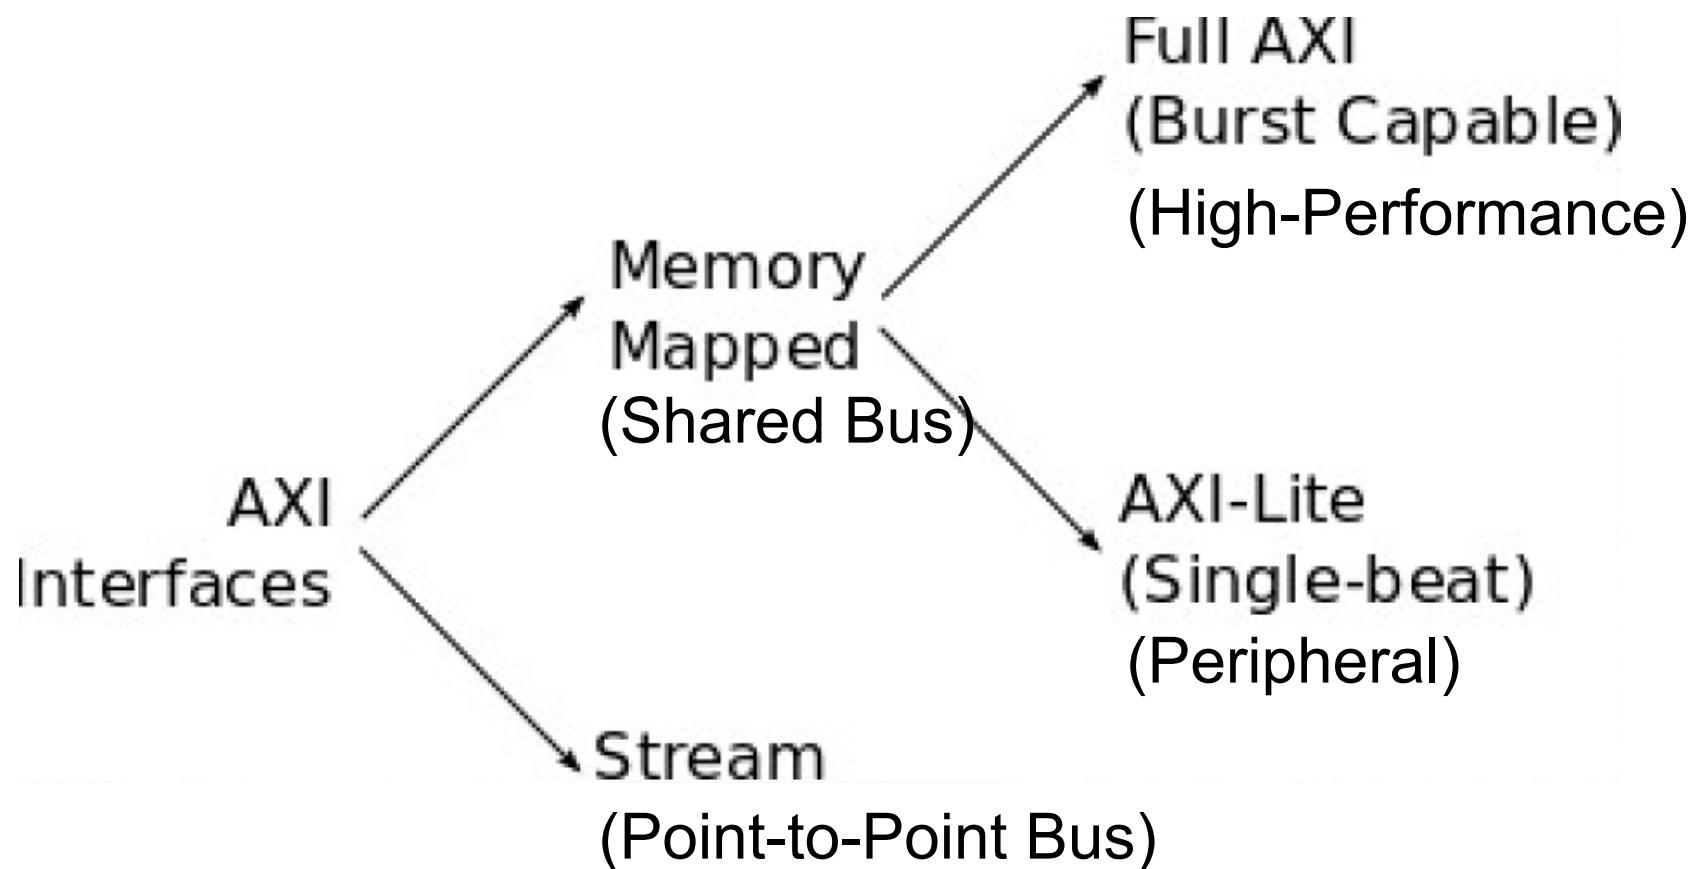
Write to Accelerator

- processor allocates buffer
- processor writes data into buffer
- ~~processor flushes cache for buffer~~
- processor initiates DMA transfer

Read from Accelerator

- processor allocates buffer
- processor initiates DMA transfer
- processor waits for DMA to complete
- ~~processor invalidates cache for buffer~~
- processor reads data from buffer

AXI Interfaces



Source: M.S. Sadri, Zynq Training

Features of AXI Interfaces

Interface	Features
MemoryMap/Full	Traditional address/data burst (single address,multiple data)
Streaming	Dataonly,burst
Lite	Traditional address/data—no burst (single address,multiple data)

Source: Building Zynq Accelerators with Vivado HLS, FPL 2013 Tutorial

Summary of AXI Full and AXI Lite Interfaces

➤ Memory mapped interfaces consist of 5 streams

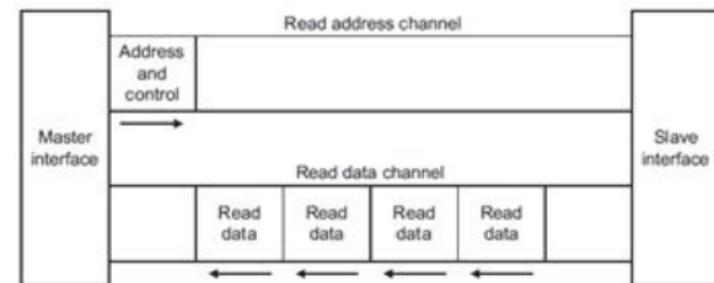
- Read Address
- Read Data
- Write Address
- Write Data
- Write Acknowledge

➤ Burst length limited to 256

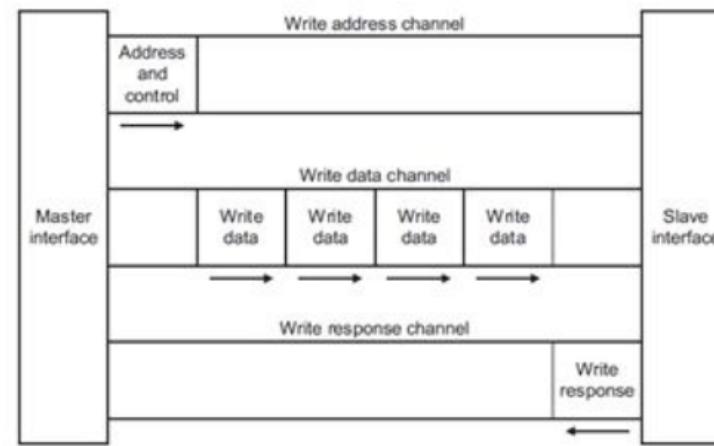
➤ Data width limited to 256 bits for Xilinx IP

➤ AXI Lite is a subset

- no bursts
- 32 bit data width only



AXI4 READ



AXI4 Write

Summary of AXI Stream Interface

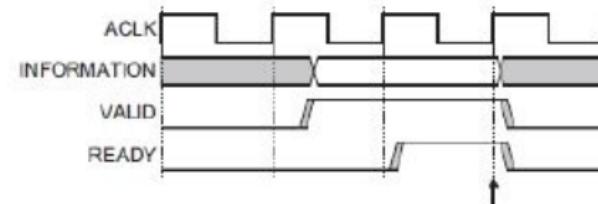
➤ AXI Streams are fully handshaked

- Data is transferred when source asserts VALID and destination asserts READY

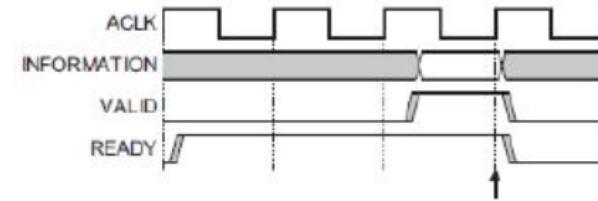
➤ ‘Information’ includes DATA and other side channel signals

- STRB
- KEEP
- LAST
- ID
- DEST
- USER

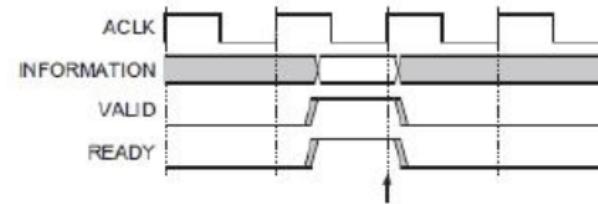
➤ Most of these are optional



Inserting Wait States

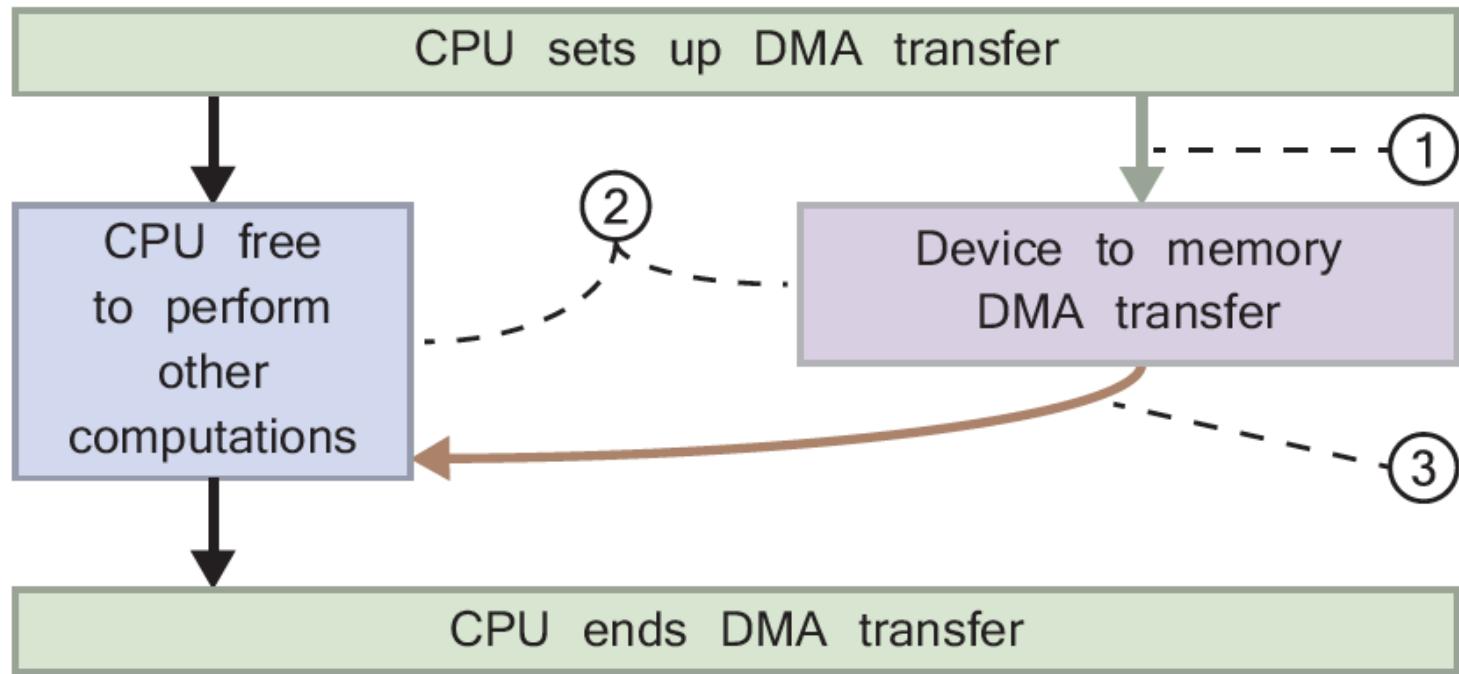


Always Ready



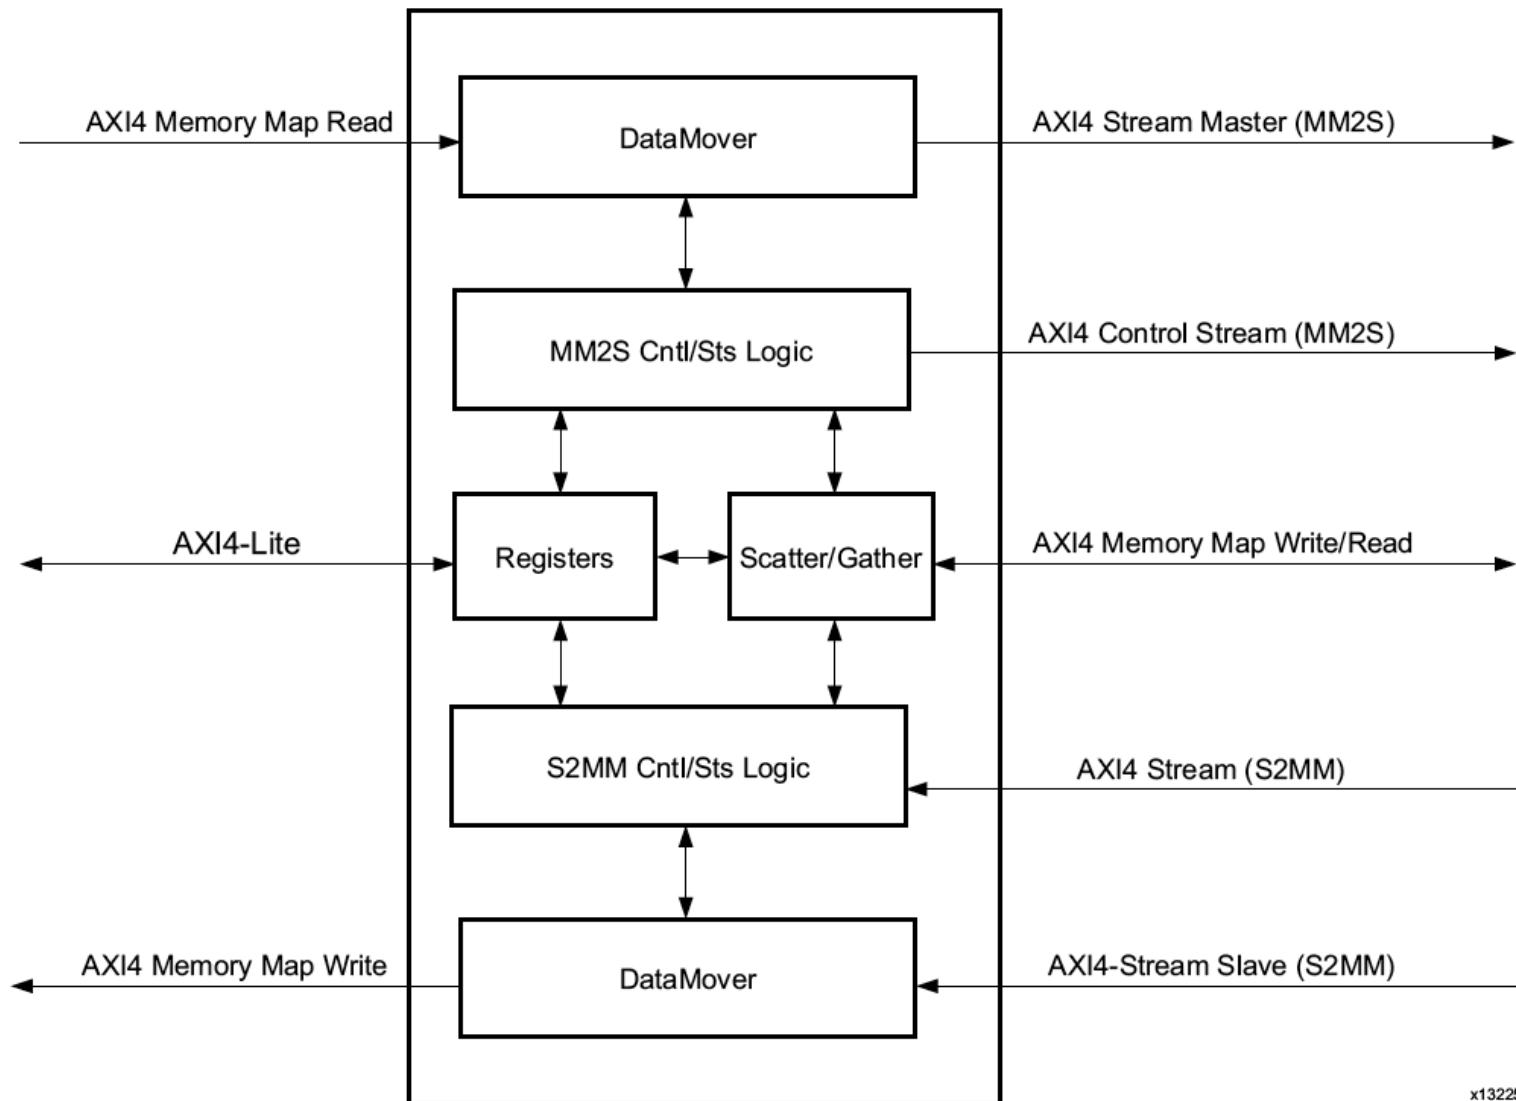
Same Cycle Acknowledge

DMA Memory Transfer Operation



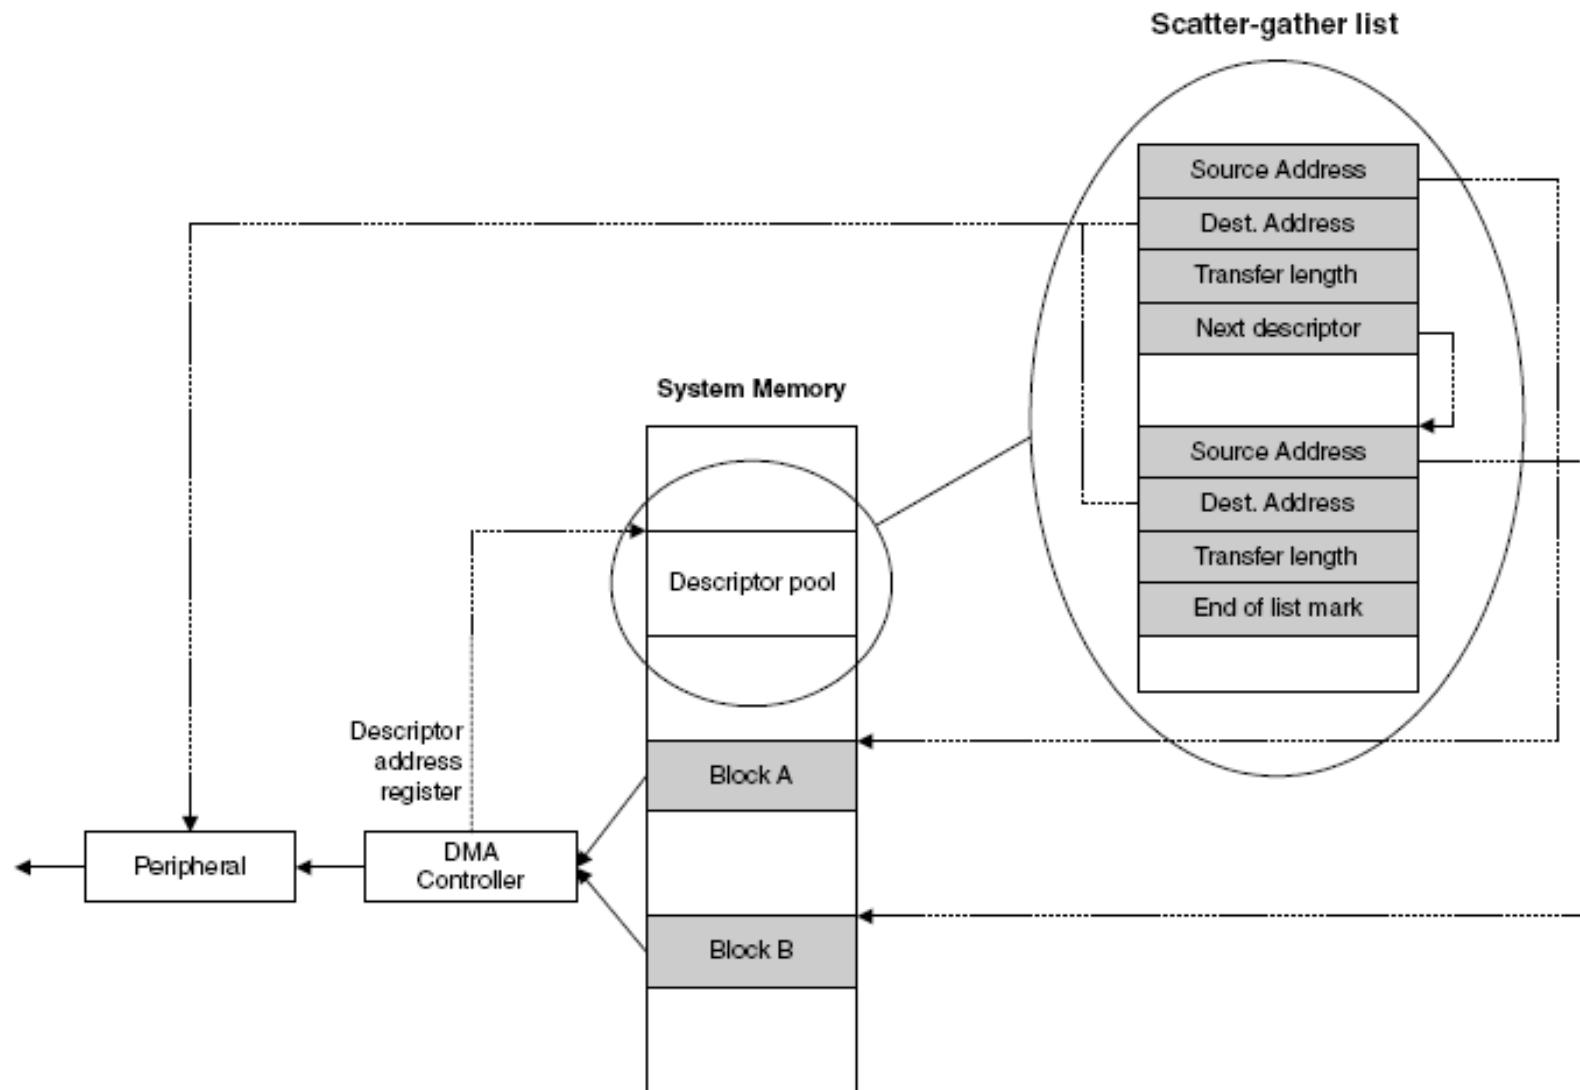
Source: The Zynq Book

Block Diagram of AXI DMA Core



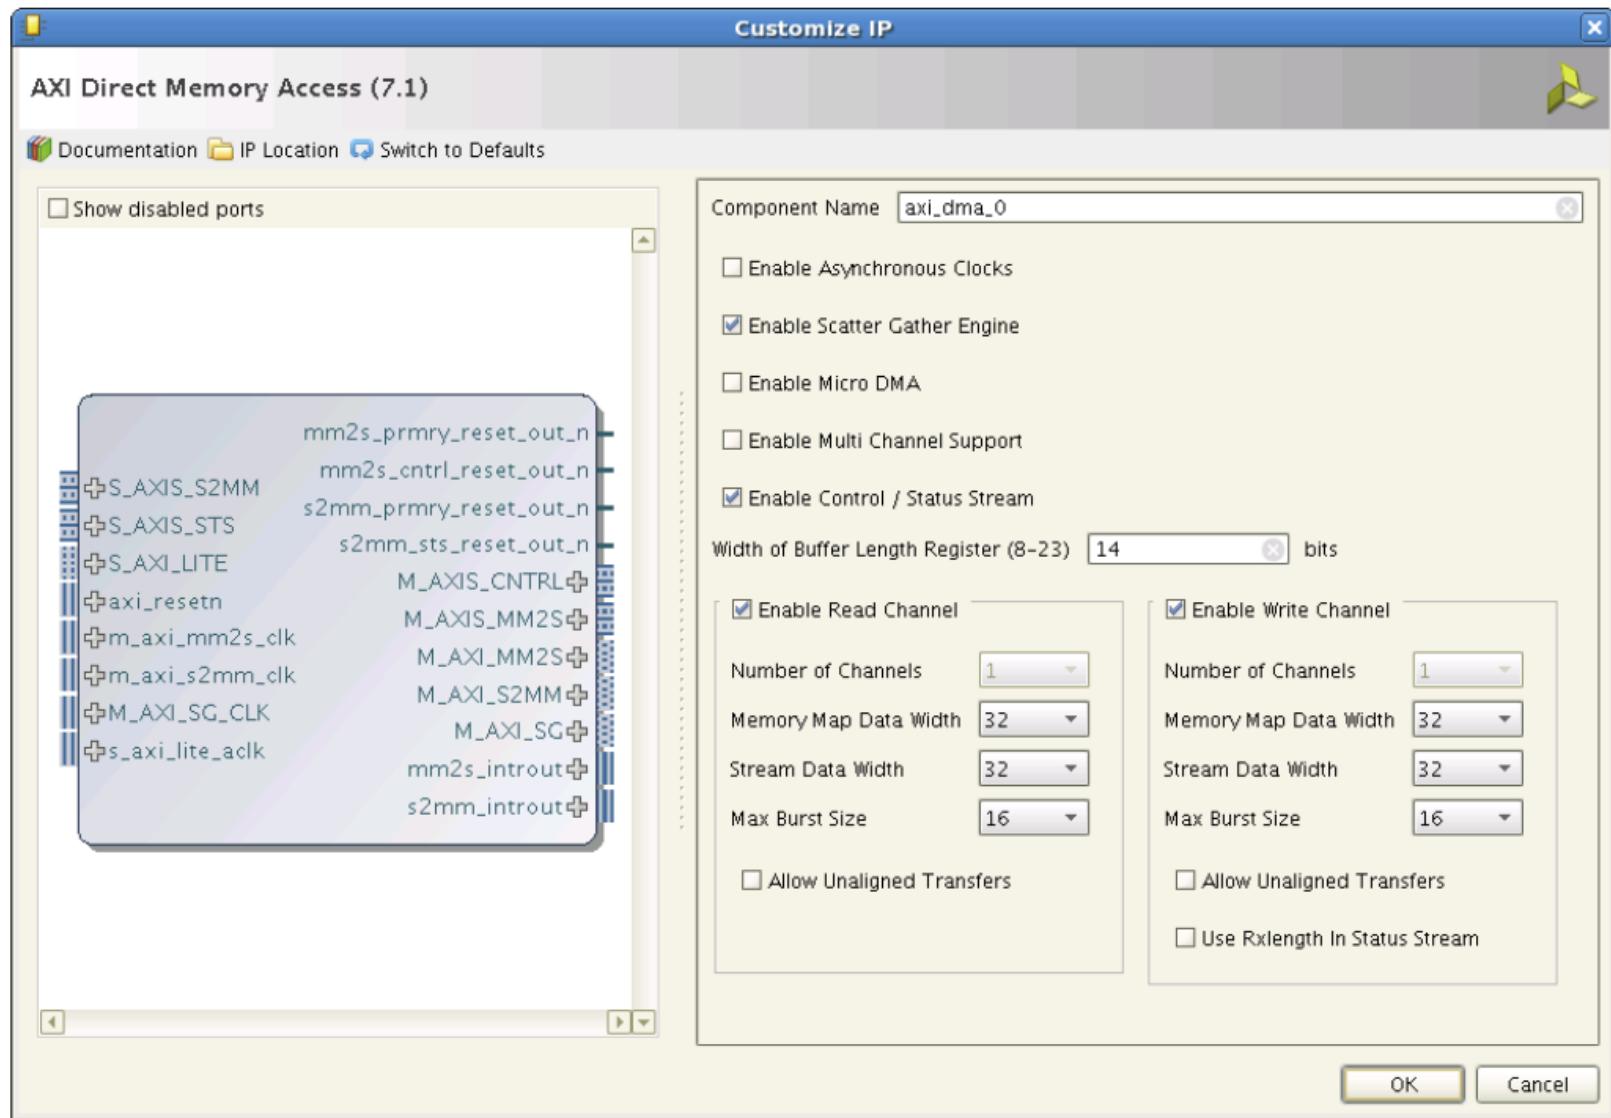
Source: LogiCORE IP AXI DMA v7.1: Product Guide

Scatter Gather DMA Mode



Source: Symbian OS Internals/13. Peripheral Support

Customizing AXI DMA Core



Source: LogiCORE IP AXI DMA v7.1: Product Guide

Parameters of AXI DMA Core (1)

Number of Channels

the number of channels: 1..16

Memory Map Data Width

data width in bits of the AXI MM2S Memory Map Read data bus: 32, 64, 128, 256, 512 or 1,024

Stream Data Width

data width in bits of the AXI MM2S AXI4-Stream Data bus:
8, 16, 32, 64, 128, 512 or 1,024;
Stream Data Width \leq Memory Map Data Width;

Max Burst Size

maximum size of burst on the AXI4-Memory Map side of MM2S: 2, 4, 8, 16, 32, 64, 128, or 256

Options of AXI DMA Core (1)

Enable Asynchronous Clocks

- 0 – all clocks inputs should be connected to the same clock signal
- 1 – separate asynchronous clocks for MM2S interface,
S2MM interface, AXI4-Lite control interface, and the
Scatter Gather Interface

Enable Scatter Gather Engine

- 0 – Simple DMA Mode operation
- 1 – Scatter Gather Mode operation; the Scatter Gather Engine included in AXI DMA

Options of AXI DMA Core (2)

Enable Micro DMA

- 0 – regular DMA
- 1 – area-optimized DMA; the maximum number of bytes per transaction = MMap_Data_width * Burst_length/8;
addressing restricted to burst boundaries

Enable Multi Channel Support

- 0 – the number of channels fixed at 1 for both directions
- 1 – the number of channels can be greater than 1

Options of AXI DMA Core (3)

Enable Control/Status Stream

0 – no AXI4 Control/Status Streams

1 – The AXI4 Control stream allows user application metadata associated with the MM2S channel to be transmitted to a target IP. The AXI4 Status stream allows user application metadata associated with the S2MM channel to be received from a target IP.

Width of Buffer Length Register (8-23)

For Simple DMA mode:

the number of valid bits in the MM2S_LENGTH and S2MM_LENGTH registers

For Scatter Gather mode:

the number of valid bits used for the Control field *buffer length* and Status field *bytes transferred* in the Scatter/Gather descriptors

Options of AXI DMA Core (4)

Allow Unaligned Transfers

Enables or disables the MM2S Data Realignment Engine (DRE). If the DRE is enabled, data reads can start from any Buffer Address byte offset, and the read data is aligned such that the first byte read is the first valid byte out on the AXI4-Stream.

Use RxLength In Status Stream

Allows AXI DMA to use a receive length field that is supplied by the S2MM target IP in the App4 field of the status packet. This gives AXI DMA a pre-determined receive byte count, allowing AXI DMA to command the exact number of bytes to be transferred.

Simple DMA Transfer Programming Sequence for MM2S channel (1)

1. Start the MM2S channel running by setting the run/stop bit to 1, MM2S_DMACR.RS = 1.
2. If desired, enable interrupts by writing a 1 to MM2S_DMACR.IOC_IrqEn and MM2S_DMACR.Err_IrqEn.
3. Write a valid source address to the MM2S_SA register.
4. Write the number of bytes to transfer in the MM2S_LENGTH register.

The MM2S_LENGTH register must be written last.

All other MM2S registers can be written in any order.

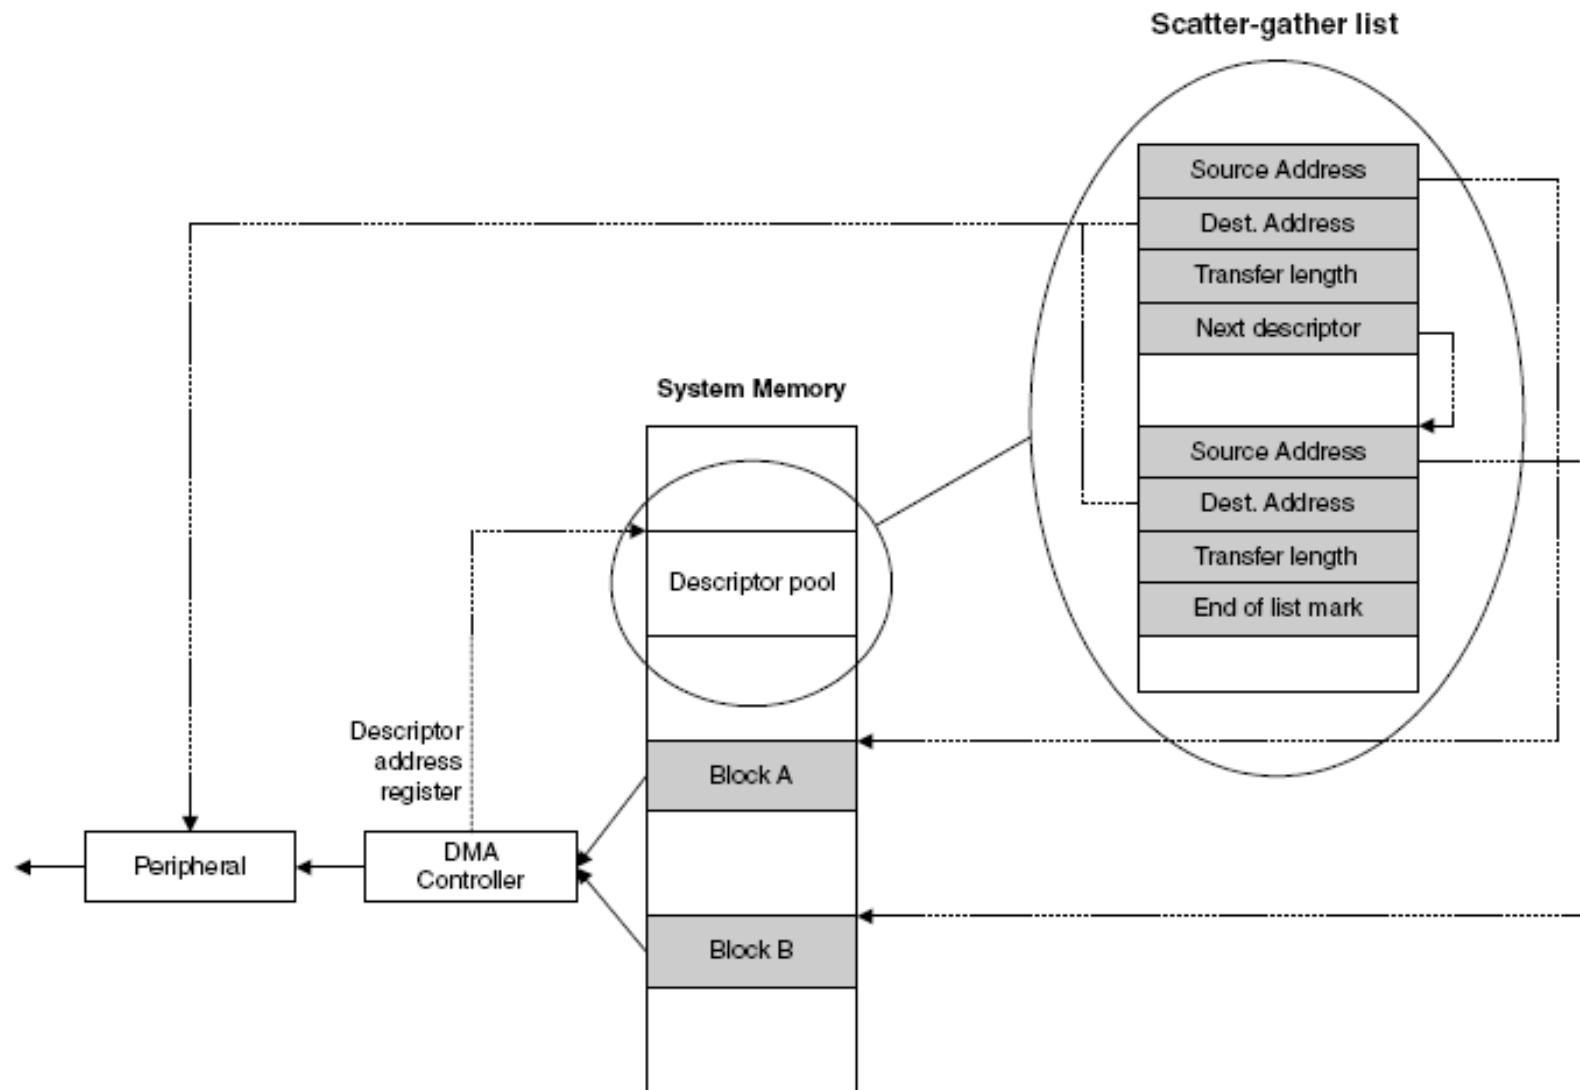
Simple DMA Transfer Programming Sequence for S2MM channel (1)

1. Start the **S2MM channel running** by setting the run/stop bit to 1, S2MM_DMACR.RS = 1.
2. If desired, **enable interrupts** by writing a 1 to S2MM_DMACR.IOC_IrqEn and S2MM_DMACR.Err_IrqEn.
3. Write a **valid destination address** to the S2MM_DA register.
4. Write the **length in bytes of the receive buffer** in the S2MM_LENGTH register.

The S2MM_LENGTH register must be written last.

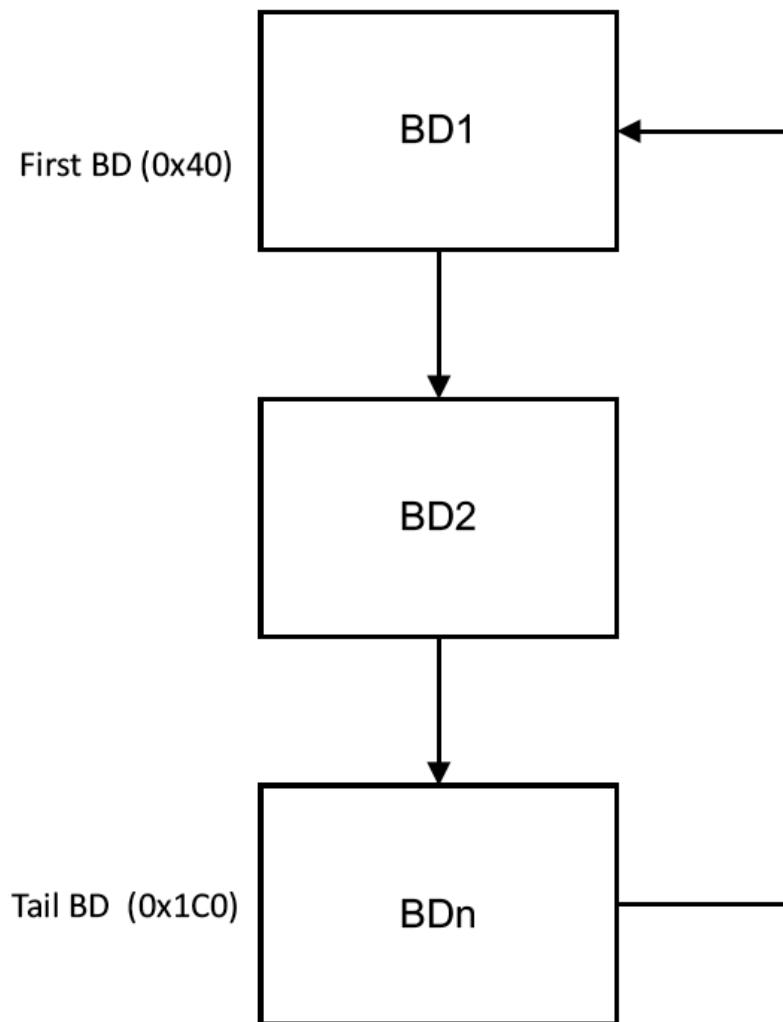
All other S2MM registers can be written in any order.

Scatter Gather DMA Mode



Source: Symbian OS Internals/13. Peripheral Support

Chain of Buffer Descriptors (BDs)



Scatter Gather DMA Transfer Programming Sequence for MM2S channel (1)

1. Write the address of the starting descriptor to the Current Descriptor register
2. Start the MM2S channel running by setting the run/stop bit to 1, MM2S_DMACR.RS = 1.
3. If desired, enable interrupts by writing a 1 to MM2S_DMACR.IOC_IrqEn and MM2S_DMACR.Err_IrqEn.
4. Write a valid address to the Tail Descriptor register.

Writing to the Tail Descriptor register triggers the DMA to start fetching the descriptors from the memory.

Simple DMA Transfer Programming Sequence for S2MM channel (1)

1. Write the address of the starting descriptor to the Current Descriptor register
2. Start the S2MM channel running by setting the run/stop bit to 1, S2MM_DMACR.RS = 1.
3. If desired, enable interrupts by writing a 1 to S2MM_DMACR.IOC_IrqEn and S2MM_DMACR.Err_IrqEn.
4. Write a valid address to the Tail Descriptor register.

Writing to the Tail Descriptor register triggers the DMA to start fetching the descriptors from the memory.